



在调试器下理解ARMv8

——虚拟内存


张银奎 2022-01-09 上海

盛格塾




IBM公司与哈佛大学合作的马克一号 (~1940)

盛格塾




马克一号的
程序纸带，两处
黑色是真正的
“补丁”，把
某些孔封上做
修改

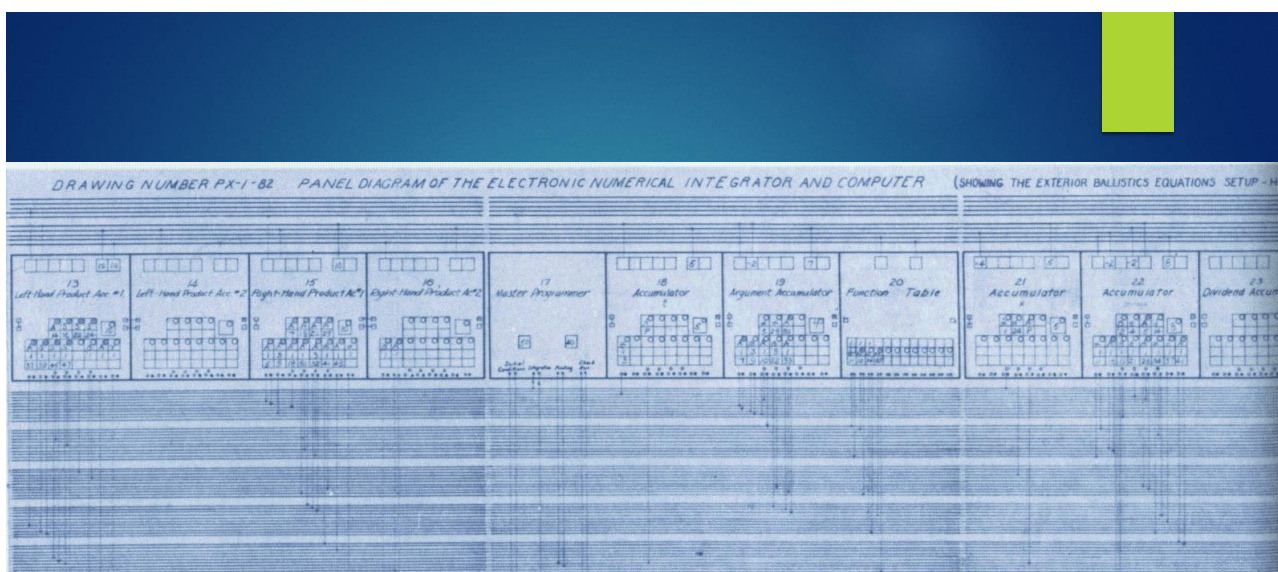
(1940 - 1943)



在戈德斯坦和埃克特的合影中，他们身后有ENIAC的四个机柜，其中左面三个便是累加器。每个累加器面板的上方有10X10的灯泡（neon）阵列，用来显示10个戴刻德的值，每一列10个灯泡，上面标有0-9十个数字，与一个戴刻德对应，点亮的灯泡代表当前值。



为ENIAC
加载程序



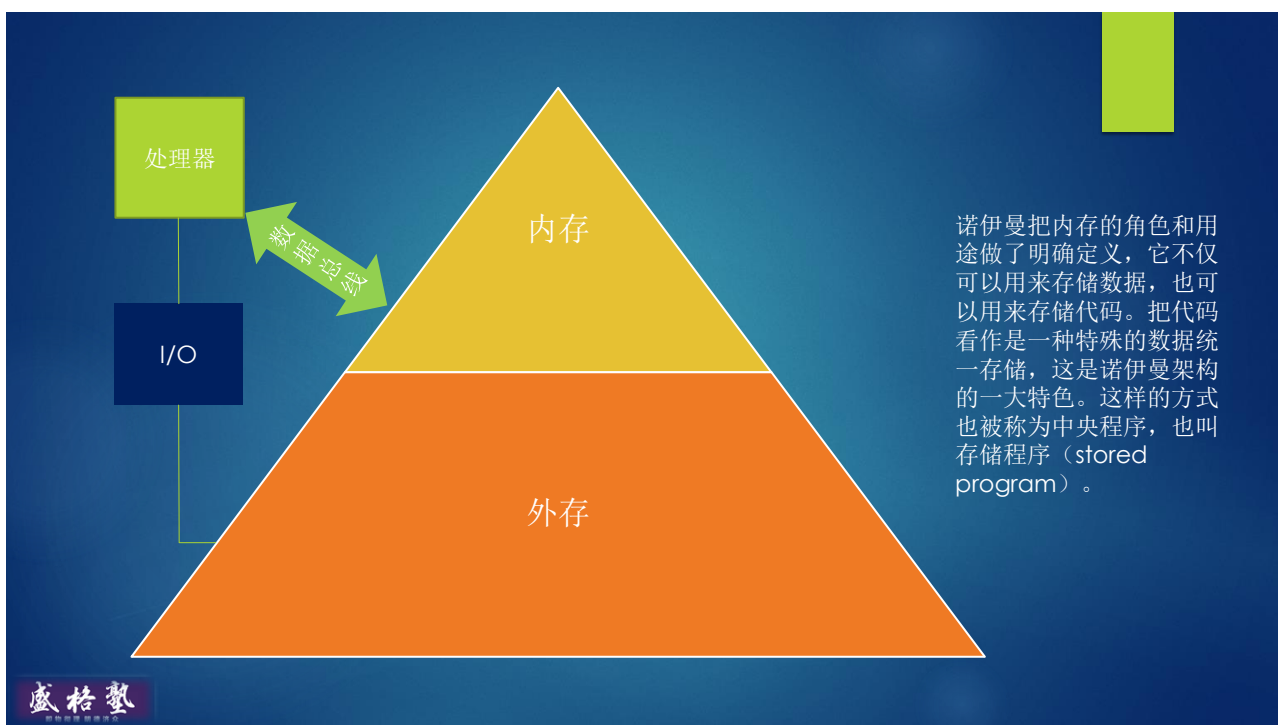
求解弹道方程的ENIAC程序

盛格黎



今天的计算机系统大多都是冯·诺伊曼架构

John von Neumann
1903-1957



Core Memory



Magnetic-core memory was the predominant form of random-access computer memory for 20 years between about 1955 and 1975. Such memory is often just called core memory, or, informally, core

威格整

VMS
(Virtual Memory System)

VAX
(Virtual Address Extension)

1975年4月，代号为Star的VAX项目开始

1975年6月，Cutler、Dick-Histvedt和Peter Lipman被任命为项目Leader，开发代号为Starlet的操作系统

威格整

VMS (Virtual Memory System)



盛格塾

Memory Management




Figure 6-6 Virtual Address Mapping into Physical Address

1979年的
PDP-11手册
中关于虚拟
内存的插图

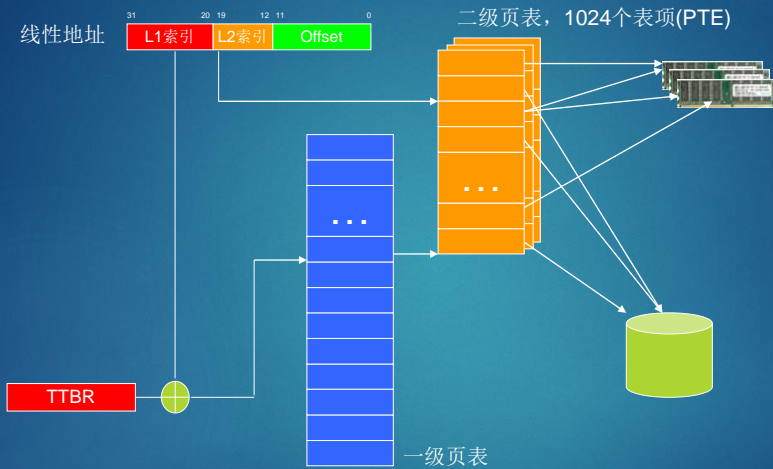
盛格塾

Page Address Register (PAR)



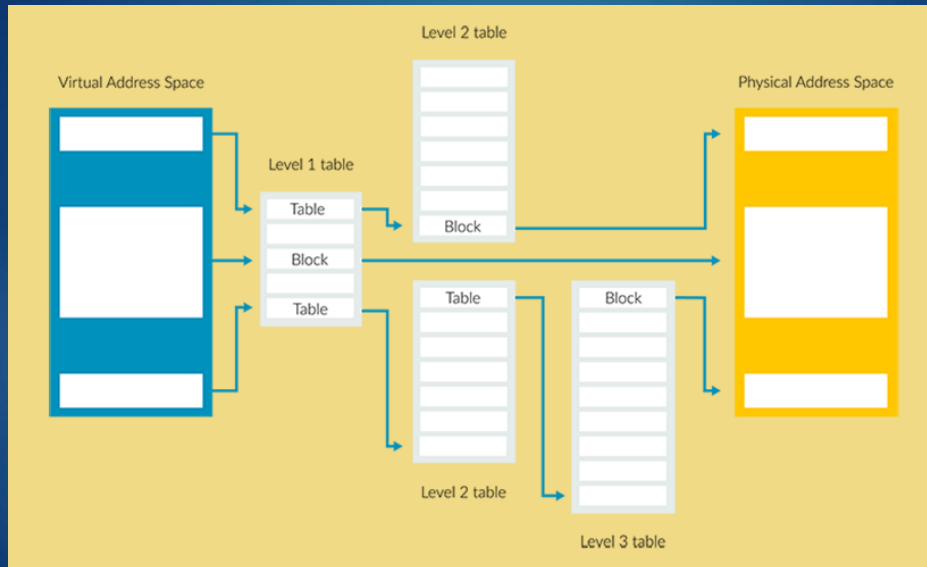
盛格塾

从线性地址到物理地址 (ARM)



- ▶ 上面是以4KB内存页(未启用PAE)为例
- ▶ $1024\text{PDE} * 1024\text{PTE} = 2^{20}$ 页

盛格塾



表项分为两种:

- 指向内存块(block)
- 指向表

10+10+12

0: kd> .formats 814bdb7e

Evaluate expression:

Hex: 814bdb7e

Decimal: -2125735042

Octal: 20122755576

Binary: 10000001 01001011 110110110110111110

Chars: .K.~

Time: ***** Invalid

Float: low -3.74427e-038 high -1.#QNAN

Double: -1.#QNAN

一级页表

```
0: kd> !dd 0c0ac000+0y1000000101*4
# c0ac814 0c656012 00000000 00000000 0c659012
# c0ac824 0c65a012 00000000 00000000 00000000
# c0ac834 0c204093 79c00493 0c205093 79400493
# c0ac844 0c65e093 0c65f093 2870b093 00000000
# c0ac854 79400493 0c0d6093 228fb093 00000000
# c0ac864 0c209093 0cbe2093 0c0cc093 7a11f093
# c0ac874 7a11e093 0cbe3093 00000000 00000000
# c0ac884 0cbd6093 00000000 0c0a4093 3c800493
```

盛格塾

二级页表

```
0: kd> !dd 0c656000+0y0010111101*4
# c6562f4 0e0bd860 0e0be860 0e0bf860 0e0c0860
# c656304 0e0c1860 0e0c2860 0e0c3860 0e0c4860
# c656314 0e0c5860 0e0c6860 0e0c7860 0e0c8860
# c656324 0e0c9092 0e0ca8c0 0e0cb092 0e0cc8c0
# c656334 0e0cd8c0 0e0ce8c0 0e0cf8c0 0e0d0082
# c656344 0e0d18c0 0e0d28c0 0e0d38c0 0e0d48c0
# c656354 0e0d58c0 0e0d68c0 0e0d78c0 0e0d88c0
# c656364 00000000 00000000 00000000 00000000
```

pfn

盛格塾

物理地址

页的基地址 + 页内偏移

```
0: kd> !db e0bd000+b7e
# e0bdb7e 48 76 69 49 73 48 79 70-65 72 76 69 73 6f 72 56 HviIsHypervisorV
# e0bdb8e 65 6e 64 6f 72 4d 69 63-72 6f 73 6f 66 74 00 48 endorMicrosoft.H
# e0bdb9e 76 69 49 73 49 6f 6d 6d-75 49 6e 55 73 65 00 48 viIsIommuInUse.H
# e0bdbae 76 6c 47 65 74 4c 70 49-6e 64 65 78 46 72 6f 6d vlGetLpIndexFrom
# e0bdbbe 41 70 69 63 49 64 00 48-76 6c 51 75 65 72 79 41 ApicId.HvIQueryA
# e0bdbce 63 74 69 76 65 48 79 70-65 72 76 69 73 6f 72 50 ctiveHypervisorP
# e0bdbde 72 6f 63 65 73 73 6f 72-43 6f 75 6e 74 00 48 76 rocessorCount.Hv
# e0bdbee 6c 51 75 65 72 79 41 63-74 69 76 65 50 72 6f 63 lQueryActiveProc
```

威格黎

```
0: kd> !db e0bdb7e
# e0bdb7e 48 76 69 49 73 48 79 70-65 72 76 69 73 6f 72 56 HviIsHypervisorV
# e0bdb8e 65 6e 64 6f 72 4d 69 63-72 6f 73 6f 66 74 00 48 endorMicrosoft.H
# e0bdb9e 76 69 49 73 49 6f 6d 6d-75 49 6e 55 73 65 00 48 viIsIommuInUse.H
# e0bdbae 76 6c 47 65 74 4c 70 49-6e 64 65 78 46 72 6f 6d vlGetLpIndexFrom
# e0bdbbe 41 70 69 63 49 64 00 48-76 6c 51 75 65 72 79 41 ApicId.HvIQueryA
# e0bdbce 63 74 69 76 65 48 79 70-65 72 76 69 73 6f 72 50 ctiveHypervisorP
# e0bdbde 72 6f 63 65 73 73 6f 72-43 6f 75 6e 74 00 48 76 rocessorCount.Hv
# e0bdbee 6c 51 75 65 72 79 41 63-74 69 76 65 50 72 6f 63 lQueryActiveProc
0: kd> db 814bdb7e
814bdb7e 48 76 69 49 73 48 79 70-65 72 76 69 73 6f 72 56 HviIsHypervisorV
814bdb8e 65 6e 64 6f 72 4d 69 63-72 6f 73 6f 66 74 00 48 endorMicrosoft.H
814bdb9e 76 69 49 73 49 6f 6d 6d-75 49 6e 55 73 65 00 48 viIsIommuInUse.H
814bdbae 76 6c 47 65 74 4c 70 49-6e 64 65 78 46 72 6f 6d vlGetLpIndexFrom
814bdbbe 41 70 69 63 49 64 00 48-76 6c 51 75 65 72 79 41 ApicId.HvIQueryA
814bdbce 63 74 69 76 65 48 79 70-65 72 76 69 73 6f 72 50 ctiveHypervisorP
814bdbde 72 6f 63 65 73 73 6f 72-43 6f 75 6e 74 00 48 76 rocessorCount.Hv
814bdbee 6c 51 75 65 72 79 41 63-74 69 76 65 50 72 6f 63 lQueryActiveProc
```

威格黎

!pte

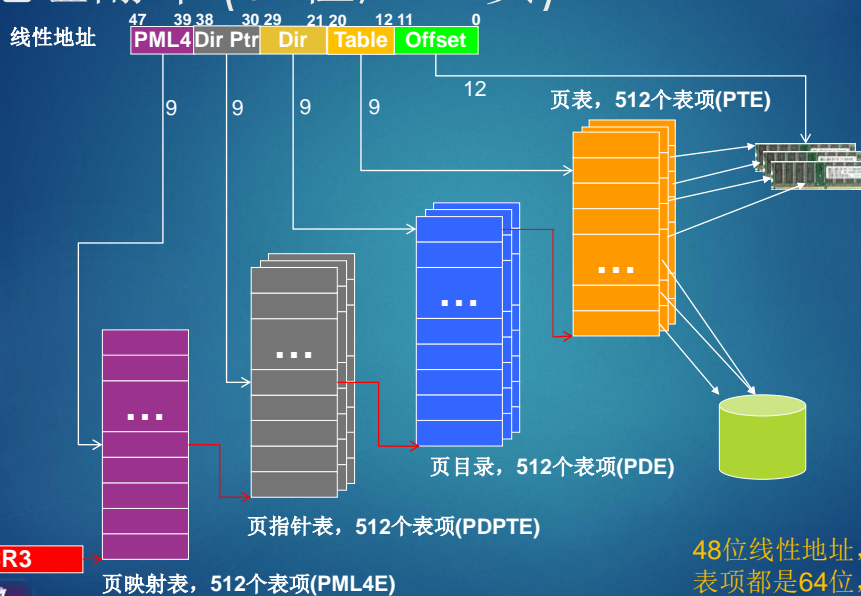
```

0: kd> !pte 814bdb7e
                VA 814bdb7e
PDE at C0300814      PTE at C02052F4
contains 0C656012    contains 0E0BD860
pfn c656  G-D-R-KA-VE not valid
                    Transition: e0bd
                    Protect: 3 - ExecuteRead

```

威格塾

地址翻译 (64位, 4KB页)



威格塾

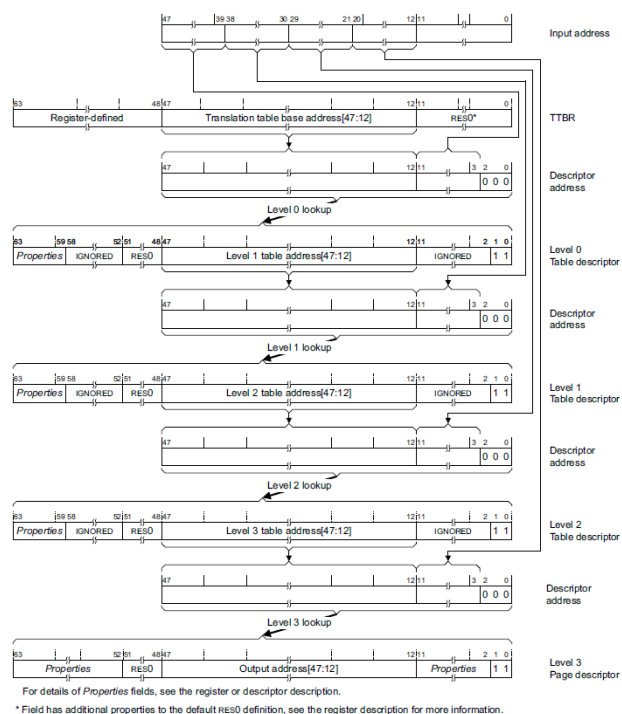
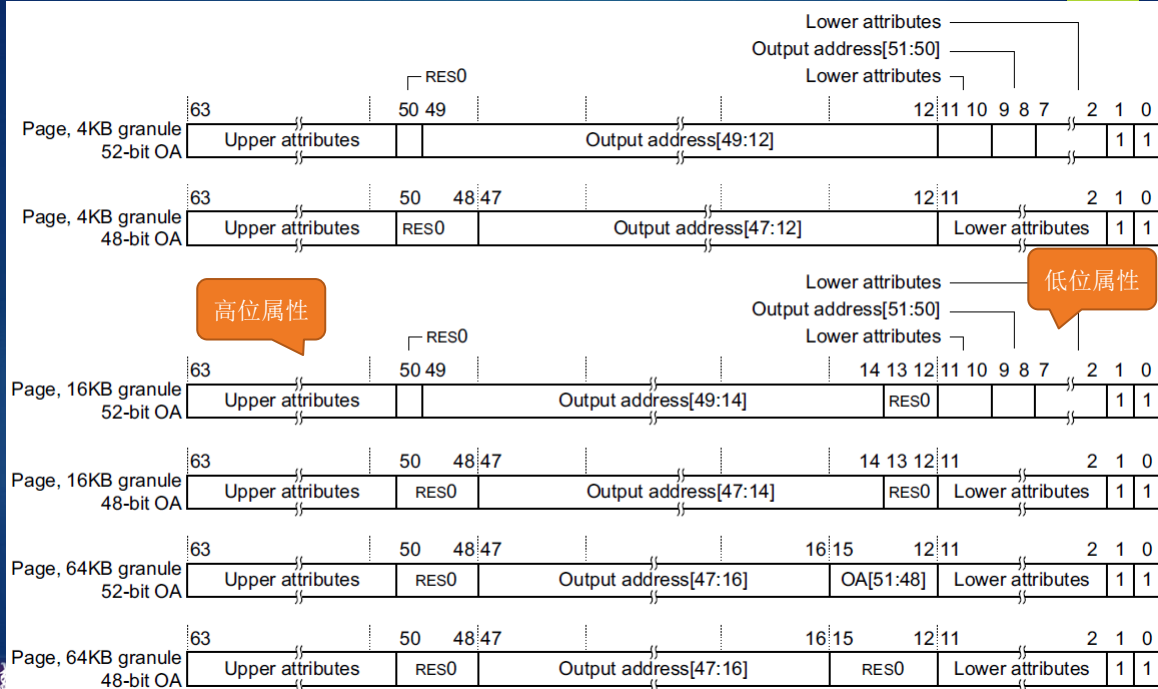


Figure K7-10 (ARMv8 ARM)

Full translation flow using the 4KB granule and starting at level 0



XN, Execute-never
Descriptor bit[54]

不要执行，
安全支持

盛格塾

Table 4-19. Format of a 4-Level Page-Table Entry that Maps a 4-KByte Page

Bit Position(s)	Contents
0 (P)	Present; must be 1 to map a 4-KByte page
1 (R/W)	Read/write; if 0, writes may not be allowed to the 4-KByte page referenced by this entry (see Section 4.6)
2 (U/S)	User/supervisor; if 0, user-mode accesses are not allowed to the 4-KByte page referenced by this entry (see Section 4.6)
3 (PWT)	Page-level write-through; indirectly determines the memory type used to access the 4-KByte page referenced by this entry (see Section 4.9.2)
4 (PCD)	Page-level cache disable; indirectly determines the memory type used to access the 4-KByte page referenced by this entry (see Section 4.9.2)
5 (A)	Accessed; indicates whether software has accessed the 4-KByte page referenced by this entry (see Section 4.8)
6 (D)	Dirty; indicates whether software has written to the 4-KByte page referenced by this entry (see Section 4.8)
7 (PAT)	Indirectly determines the memory type used to access the 4-KByte page referenced by this entry (see Section 4.9.2)
8 (G)	Global; if CR4.PGE = 1, determines whether the translation is global (see Section 4.10); ignored otherwise
11:9	Ignored
(M-1):12	Physical address of the 4-KByte page referenced by this entry
51:M	Reserved (must be 0)
58:52	Ignored
62:59	Protection key; if CR4.PKE = 1, determines the protection key of the page (see Section 4.6.2); ignored otherwise
63 (XD)	If IA32_EFER.NXE = 1, execute-disable (if 1, instruction fetches are not allowed from the 4-KByte page controlled by this entry; see Section 4.6); otherwise, reserved (must be 0)

Execute-disable bit
enable (NXE)

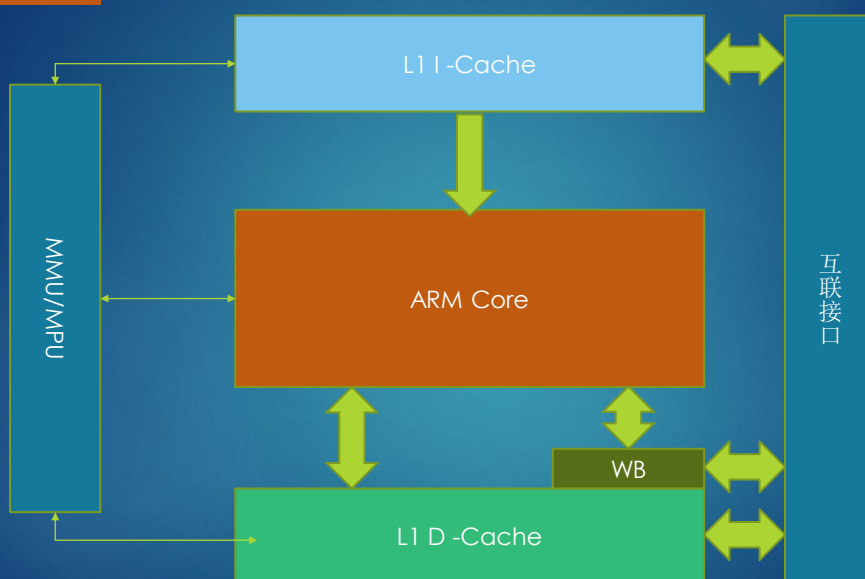
老大哥
叫XD和
NX

盛格塾



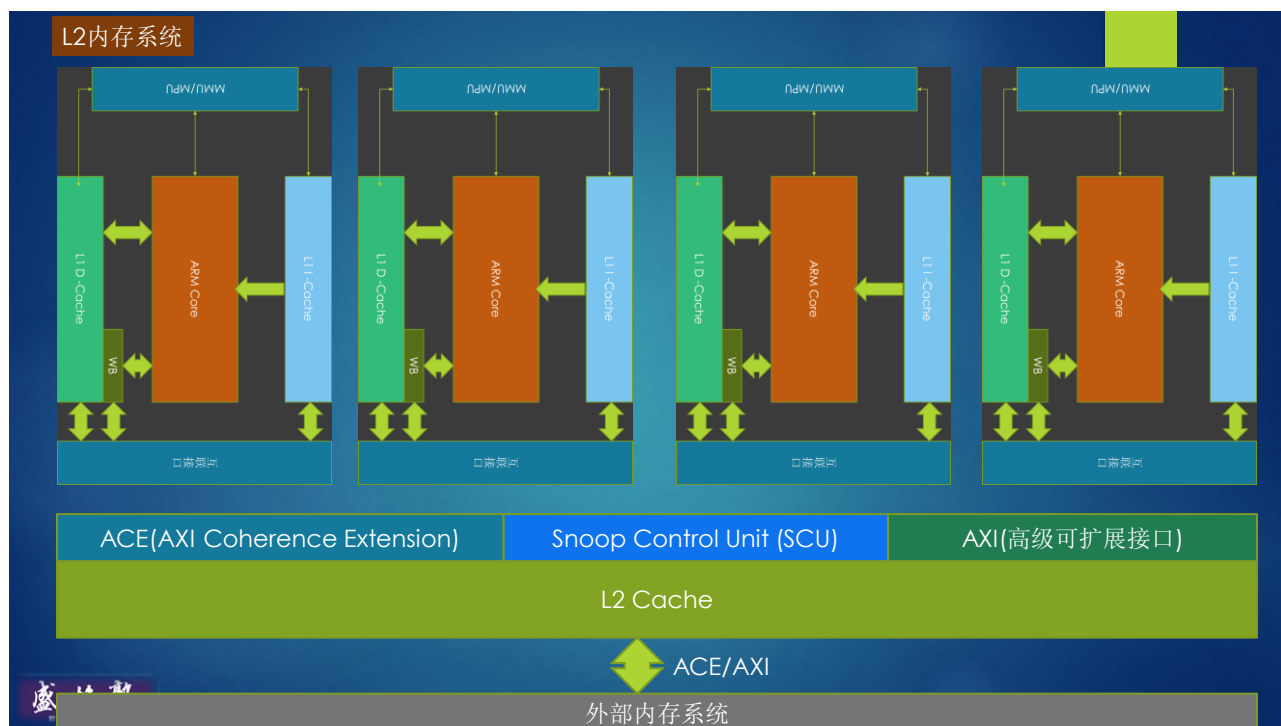
盛格塾

L1内存系统



MMU = Memory Management Unit
MPU = Memory Protection Unit

盛格塾



- ✓ 5: Memory Management Unit
 - 5.1 About the MMU
 - > 5.2 TLB organization
 - 5.3 TLB match process
 - > 5.4 External aborts
- ✓ 6: Level 1 Memory System
 - 6.1 About the L1 memory system
 - > 6.2 Cache behavior
 - 6.3 Support for v8 memory types
 - > 6.4 L1 Instruction memory system
 - > 6.5 L1 Data memory system
 - > 6.6 Data prefetching
 - > 6.7 Direct access to internal memory
- ✓ 7: Level 2 Memory System
 - 7.1 About the L2 memory system
 - > 7.2 Snoop Control Unit
 - > 7.3 ACE master interface
 - > 7.4 CHI master interface
 - 7.5 Additional memory attributes

ARM® Cortex®-A53 MPCore Processor
Revision: r0p2
Technical Reference Manual

如何找到页表
呢？

TTBR

Translation Table Base Register

盛格整

32位TTBR

Bits	Name	Function
[31:x]	BADDR	Translation table base address, bits[31:x].
[x-1:7]	-	Reserved.
[6]	IRGN[0]	Inner region bit 0.
[5]	NOS	Not Outer Shareable bit.
[4:3]	RGN	Region bits.
[2]	-	Reserved.
[1]	S	Shareable bit.
[0]	C/IRGN[1]	Cacheable bit. / Inner region bit 1.

盛格整

64位TTBR



Bits	Name	Function
[63:48]	ASID	An ASID for the translation table base address.
[47:1]	BADDR	Translation table base address, bits[39:x].
[0]	-	Reserved.

威格黎

VA

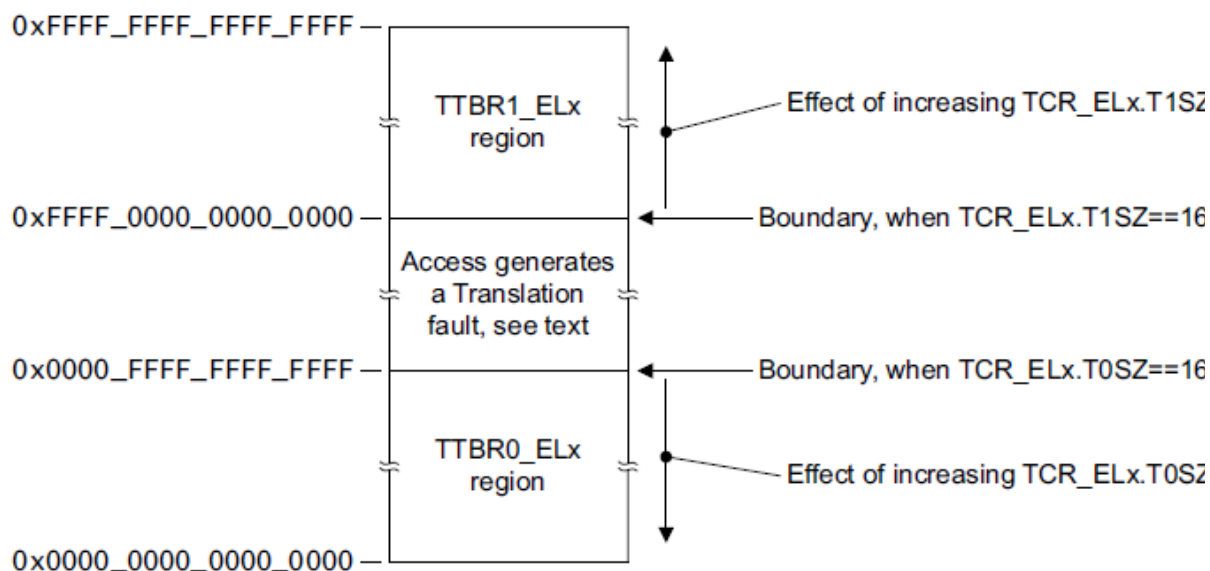




Figure D4-14 AArch64 TTBR_n boundaries and VA range

TTBRO和TTBR1

TTBRO用于翻译用户空间
TTBR1用于翻译内核空间



TTBRO



TTBR1



格物

```
[10474.137272] pgd_val = 0xd4c65003
[10474.137277] pgd_index = 511
[10474.137282] pud_val = 0xd4c65003
[10474.137287] pmd_val = 0xdf62e003
[10474.137292] pmd_index = 433
[10474.137297] pte_val = 0xe80000d4931f53
[10474.137302] pte_index = 57
[10474.137307] page_addr = e80000d4931000, page_offset = 638
[10474.137313] vaddr = 7ff6239638, paddr = e80000d4931638
[10474.137319] physical address is e80000d4931638. va 0000007ff6239638
```

通过llaolao驱动打印出的页表描述符





盛格塾

```

/*
 * vm_fault is filled by the the pagefault handler and passed to the vma's
 * ->fault function. The vma's ->fault is responsible for returning a bitmask
 * of VM_FAULT_xxx flags that give details about how the fault was handled.
 *
 * pgoff should be used in favour of virtual_address, if possible.
 */
struct vm_fault {
    unsigned int flags;          /* FAULT_FLAG_xxx flags */
    pgoff_t pgoff;              /* Logical page offset based on vma */
    void __user *virtual_address; /* Faulting virtual address */

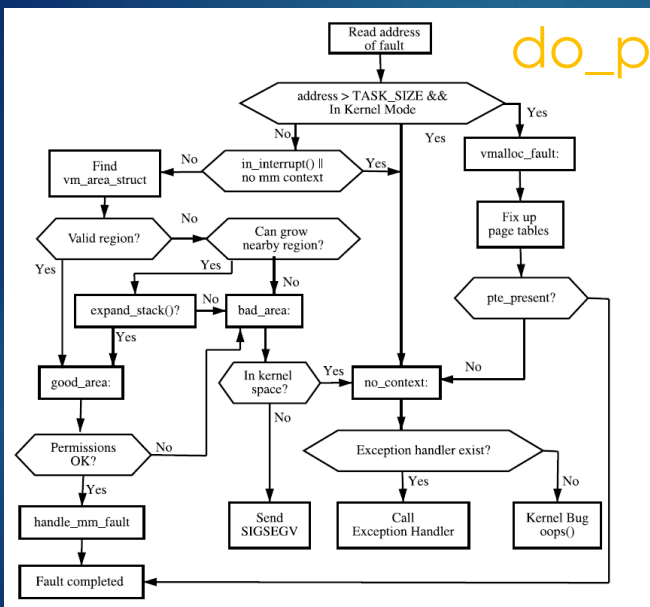
    struct page *cow_page;      /* Handler may choose to COW */
    struct page *page;          /* ->fault handlers should return a
     * page here, unless VM_FAULT_NOPAGE
     * is set (which is also implied by
     * VM_FAULT_ERROR).
     */
    /* for ->map_pages() only */
    pgoff_t max_pgoff;          /* map pages for offset from pgoff till
     * max_pgoff inclusive */
    pte_t *pte;                /* pte entry associated with ->pgoff */

```

<linux/mm.h>

盛格塾

do_page_fault()



盛格塾

细分

Minor PF

- 不需要从磁盘换进页
- 例如：访问没有分配过的malloc内存块

Major PF

- 需要从磁盘读取内容
- 情况1：换进交换出去的页
- 情况2：访问文件映射的页

盛格塾


Fields Management for window `l:Def`, whose current sort field is %CPU
 Navigate with Up/Dn, Right selects for move then <Enter> or Left commits,
 'd' or <Space> toggles display, 's' sets sort. Use 'q' or <Esc> to end!

```
* PID      = Process Id          TIME    = CPU Time
* USER     = Effective User Name SWAP     = Swapped Size (KiB)
* PR       = Priority            CODE     = Code Size (KiB)
* NI       = Nice Value          DATA    = Data+Stack (KiB)
* VIRT     = Virtual Image (KiB) nMaj     = Major Page Faults
* RES      = Resident Size (KiB) nMin    = Minor Page Faults
* SHR      = Shared Memory (KiB) nDRT    = Dirty Pages Count
* S         = Process Status      WCHAN   = Sleeping in Function
* %CPU     = CPU Usage            Flags   = Task Flags <sched.h>
* %MEM     = Memory Usage (RES)   CGROUPS = Control Groups
* TIME+    = CPU Time, hundredths SUPGIDS = Supp Groups IDs
* COMMAND  = Command Name/Line   SUPGRPS = Supp Groups Names
* PPID     = Parent Process pid   TGID    = Thread Group Id
* UID      = Effective User Id    ENVIRON  = Environment vars
* RUID     = Real User Id         vMj     = Major Faults delta
* RUSER    = Real User Name       vMn     = Minor Faults delta
* SUID     = Saved User Id        USED     = Res+Swap Size (KiB)
* SUSER    = Saved User Name      nsIPC    = IPC namespace Inode
* GID      = Group Id            nsMNT    = MNT namespace Inode
* GROUP    = Group Name          nsNET    = NET namespace Inode
* PGRP     = Process Group Id     nsPID    = PID namespace Inode
* TTY      = Controlling Tty      nsUSER   = USER namespace Inode
* TPGID    = Tty Process Grp Id   nsUTS    = UTS namespace Inode
* SID      = Session Id
* nTH      = Number of Threads
* P         = Last Used Cpu (SMP)
```

需要较高版本的
TOP命令(WSL的
可以)

PAR_EL1

► Physical Address Register, EL1

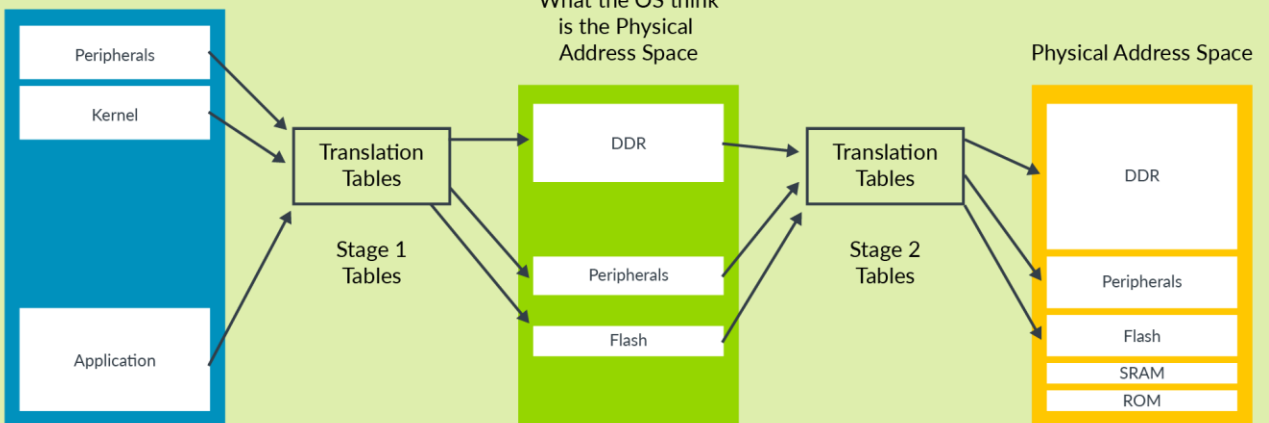


CPUMERRSR_EL1

- ▶ **CPU Memory Error Syndrome Register, EL1**
- ▶ Holds the number of memory errors that have occurred in the following L1 and L2 RAMs:
 - ▶ • L1-I Tag RAM.
 - ▶ • L1-I Data RAM.
 - ▶ • L1-D Tag RAM.
 - ▶ • L1-D Data RAM.
 - ▶ • L2 TLB RAM.

盛格塾

Virtual Address Space



盛格塾

<https://developer.arm.com/documentation/102142/0100/Stage-2-translation>

G8.2.77 HTTBR, Hyp Translation Table Base Register

The HTTBR characteristics are:

Purpose

Holds the base address of the translation table for the initial lookup for stage 1 of an address translation in the EL2 translation regime, and other information for this translation regime.

Configurations

AArch32 System register HTTBR bits [47:1] are architecturally mapped to AArch64 System register **TTBR0_EL2**[47:1].

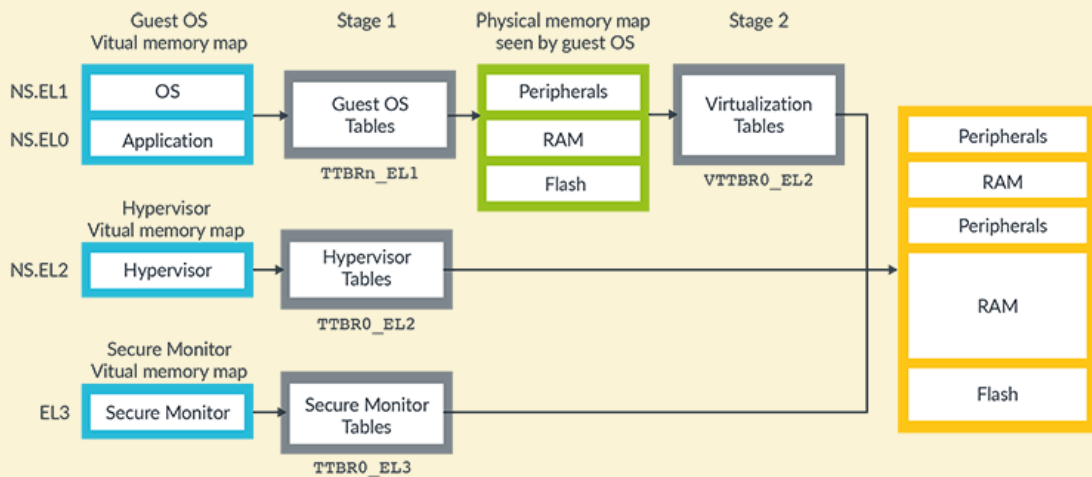
This register is present only when AArch32 is supported at EL0. Otherwise, direct accesses to HTTBR are UNDEFINED.

If EL2 is not implemented, this register is RES0 from EL3.

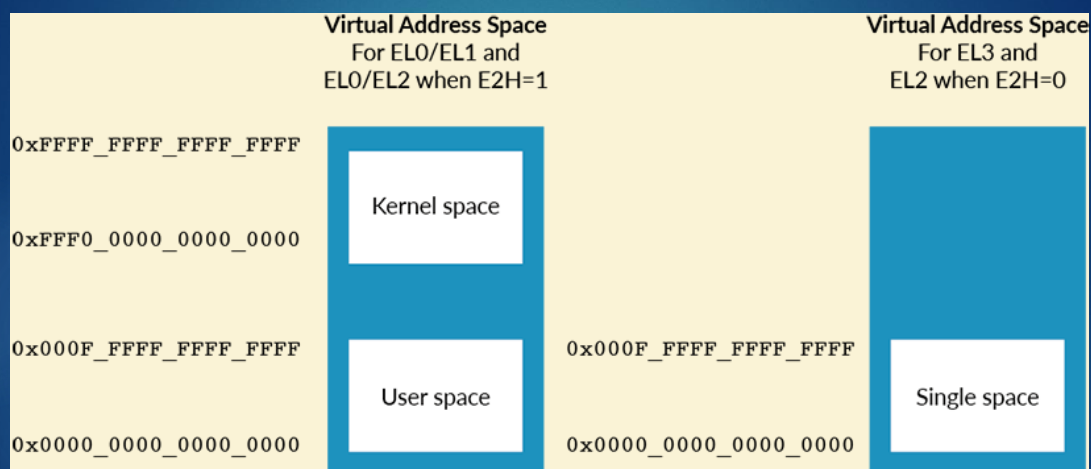
Attributes

HTTBR is a 64-bit register.


威格整



威格整



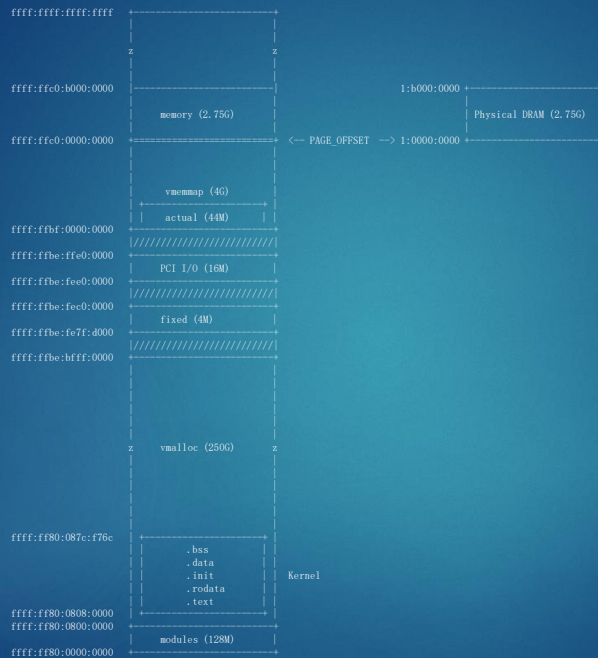
威格黎



Arch64地址空间布局

CONFIG_ARM64_PAGE_SHIFT=12
CONFIG_PGTABLE_LEVELS=3
CONFIG_ARM64_VA_BITS=39

威格黎



Arch64内核空间布局

[0.000000] Virtual kernel memory layout:

```
modules : 0xffffffff8000000000 - 0xffffffff8008000000 ( 128 MB)
vmalloc : 0xffffffff8008000000 - 0xffffffffbdbc000000 ( 246 GB)
 .init : 0xffffffff8008e30000 - 0xffffffff8008f30000 ( 1024 KB)
 .text : 0xffffffff8008080000 - 0xffffffff8008a40000 ( 9984 KB)
 .rodata : 0xffffffff8008a40000 - 0xffffffff8008e30000 ( 4032 KB)
 .data : 0xffffffff8008f30000 - 0xffffffff80090a9008 ( 1509 KB)
vmemmap : 0xffffffffbdc0000000 - 0xffffffffbfc0000000 ( 8 GB maximum)
          0xffffffffbdc0008000 - 0xffffffffbdc3fc0000 ( 63 MB actual)
fixed : 0xffffffffbffe7fb000 - 0xffffffffbffe000000 ( 4116 KB)
PCI I/O : 0xffffffffbffe000000 - 0xffffffffbffe000000 ( 16 MB)
memory : 0xffffffc000200000 - 0xffffffc0ff000000 ( 4078 MB)
```

内核启动时打印出的地址空间布局

```

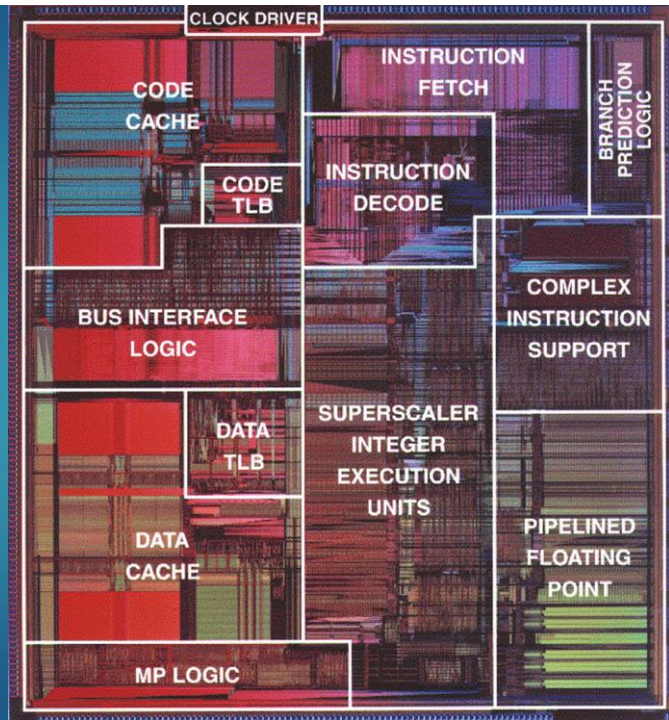
1  =====
2  Memory Layout on AArch64 Linux
3  =====
4
5  Author: Catalin Marinas <catalin.marinas@arm.com>
6
7  This document describes the virtual memory layout used by the AArch64
8  Linux kernel. The architecture allows up to 4 levels of translation
9  tables with a 4KB page size and up to 3 levels with a 64KB page size.
10
11 AArch64 Linux uses either 3 levels or 4 levels of translation tables
12 with the 4KB page configuration, allowing 39-bit (512GB) or 48-bit
13 (256TB) virtual addresses, respectively, for both user and kernel. With
14 64KB pages, only 2 levels of translation tables, allowing 42-bit (4TB)
15 virtual address, are used but the memory layout is the same.
16
17 ARMv8.2 adds optional support for Large Virtual Address space. This is
18 only available when running with a 64KB page size and expands the
19 number of descriptors in the first level of translation.
20
21 User addresses have bits 63:48 set to 0 while the kernel addresses have
22 the same bits set to 1. TTBRx selection is given by bit 63 of the
23 virtual address. The swapper_pg_dir contains only kernel (global)
24 mappings while the user pgd contains only user (non-global) mappings.
25 The swapper_pg_dir address is written to TTBR1 and never written to
26 TTBR0.
27
28
29 AArch64 Linux memory layout with 4KB pages + 4 levels (48-bit)::

```

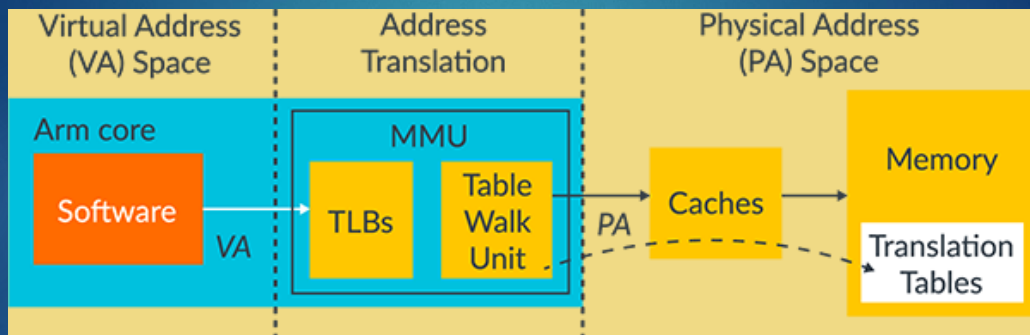


TLB

- ▶ Translation Lookaside Buffer
- ▶ a small associative memory that caches virtual to physical page table resolutions



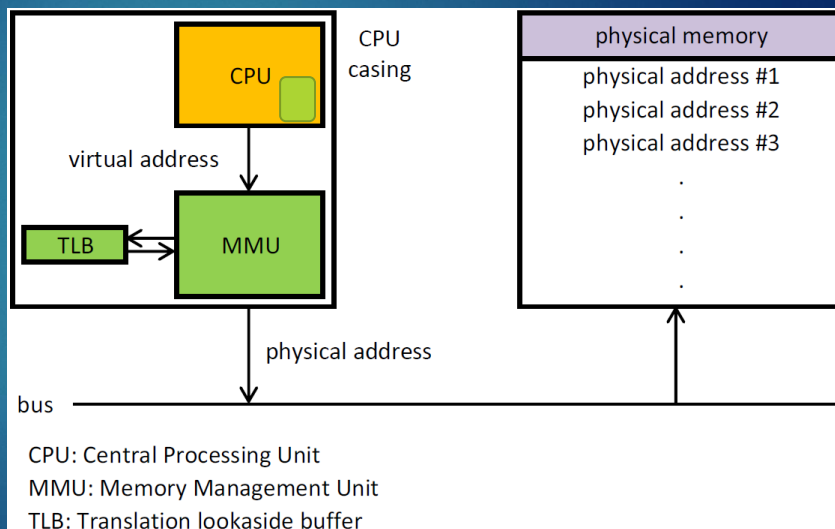
威格黎



威格黎

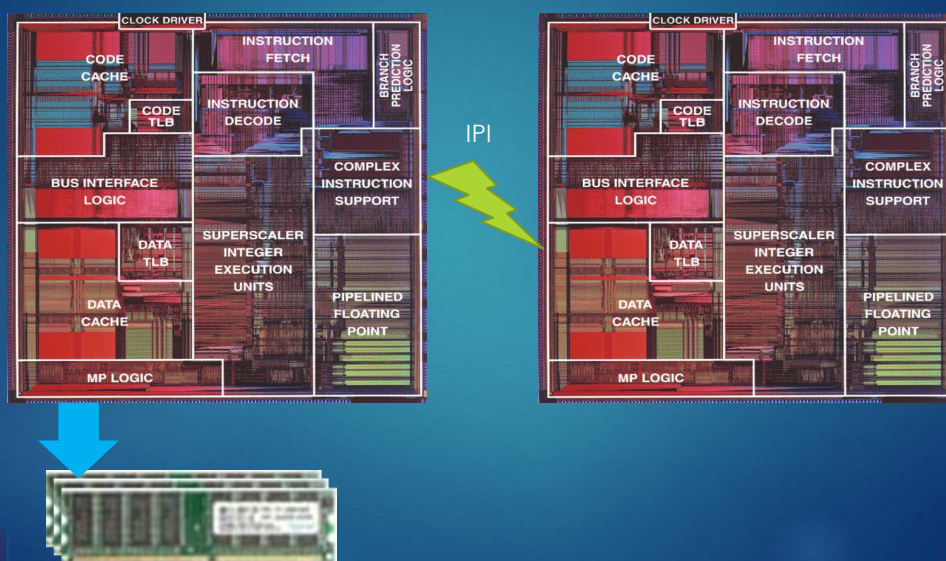
MMU

- ▶ **Memory Management Unit (MMU):** Device mapping logical (virtual) addresses to physical addresses
- ▶ **Logical address** – process view of memory
- ▶ **Physical address** – MMU view of memory





威格黎

TLB同步



威格黎



多核趋势让TLB同步的挑战更大

威格黎

```
Jan 30 22:14:43 darkstar kernel: watchdog: BUG: soft lockup - CPU#2 stuck for 44s! [worker:131042]
Jan 30 22:14:43 darkstar kernel: Modules linked in: cdc_acm rpcsec_gss_krb5 dm_mod vhost_net tun
vhost macvtap tap macvlan bonding xt_MASQUERADE iptable_nat nf_nat nf_conntrack nf_defrag_ip>
Jan 30 22:14:43 darkstar kernel: vfio_pci irqbypass vfio_virqfd vfio_iommu_type1 vfio
Jan 30 22:14:43 darkstar kernel: CPU: 2 PID: 131042 Comm: worker Tainted: G      L 5.4.12-arch1-1 #1
Jan 30 22:14:43 darkstar kernel: Hardware name: Gateway GT350 F1/GT350 F1, BIOS P03 07/26/2010
Jan 30 22:14:43 darkstar kernel: RIP: 0010:smp_call_function_many+0x21d/0x280
Jan 30 22:14:43 darkstar kernel: Code: e8 88 c8 7c 00 3b 05 d6 c1 20 01 89 c7 0f 83 7a fe ff 48 63 c7 48
8b 0b 48 03 0c c5 20 f9 f7 9d 8b 41 18 a8 01 74 0a f3 90 <8b> 51 18 83 e2 01 75 f>
Jan 30 22:14:43 darkstar kernel: RSP: 0018:ffffbc648b4ffc8 EFLAGS: 00000202 ORIG_RAX: ffffffff13
Jan 30 22:14:43 darkstar kernel: RAX: 0000000000000003 RBX: ffff9a7c5f8aba40 RCX: ffff9a7c5f8312c0
Jan 30 22:14:43 darkstar kernel: RDX: 0000000000000001 RSI: 0000000000000000 RDI: 0000000000000000
Jan 30 22:14:43 darkstar kernel: RBP: ffffffff9ce7ee90 R08: ffff9a7c5f8aba48 R09: 0000000000000006
Jan 30 22:14:43 darkstar kernel: R10: ffff9a7c5f8aba48 R11: 0000000000000005 R12: ffff9a7c5f8aa340
Jan 30 22:14:43 darkstar kernel: R13: ffff9a7c5f8aba48 R14: 0000000000000001 R15: 0000000000000140
Jan 30 22:14:43 darkstar kernel: FS: 00007f58b0dbf700(0000) GS:ffff9a7c5f880000(0000)
knlGS:0000000000000000
Jan 30 22:14:43 darkstar kernel: CS: 0010 DS: 0000 ES: 0000 CR0: 0000000080050033
Jan 30 22:14:43 darkstar kernel: CR2: 0000000000000030 CR3: 00000000808eb8004 CR4: 00000000000226e0
Jan 30 22:14:43 darkstar kernel: DR0: 000000007ffe0270 DR1: 00000000002f2040 DR2: 0000000000000000
Jan 30 22:14:43 darkstar kernel: DR3: 0000000000000000 DR6: 00000000ffff0ff0 DR7: 0000000000000400
```

<https://obs.archlinux.org/viewtopic.php?id=252523>

威格黎

```

Jan 30 22:14:43 darkstar kernel: Call Trace:
Jan 30 22:14:43 darkstar kernel: flush_tlb_mm_range+0xed/0x150
Jan 30 22:14:43 darkstar kernel: tlb_flush_mmu+0xa4/0x160
Jan 30 22:14:43 darkstar kernel: tlb_finish_mmu+0x3d/0x70
Jan 30 22:14:43 darkstar kernel: unmap_region+0xf4/0x130
Jan 30 22:14:43 darkstar kernel: __do_munmap+0x255/0x4c0
Jan 30 22:14:43 darkstar kernel: __vm_munmap+0x67/0xb0
Jan 30 22:14:43 darkstar kernel: __x64_sys_munmap+0x28/0x30
Jan 30 22:14:43 darkstar kernel: do_syscall_64+0x4e/0x140
Jan 30 22:14:43 darkstar kernel: entry_SYSCALL_64_after_hwframe+0x44/0xa9
Jan 30 22:14:43 darkstar kernel: RIP: 0033:0x7f5c028d10db
Jan 30 22:14:43 darkstar kernel: Code: 8b 15 a9 5d 0c 00 f7 d8 64 89 02 48 c7 c0 ff ff ff ff eb 89 66 2e 0f 1f
84 00 00 00 00 90 f3 0f 1e fa b8 0b 00 00 00 0f 05 <48> 3d 01 f0 ff ff 73 0>
Jan 30 22:14:43 darkstar kernel: RSP: 002b:00007f58b0dbd038 EFLAGS: 00000206 ORIG_RAX:
000000000000000b
Jan 30 22:14:43 darkstar kernel: RAX: ffffffffda RBX: 00007f58da31e9c0 RCX: 00007f5c028d10db
Jan 30 22:14:43 darkstar kernel: RDX: 0000000000000000 RSI: 000000000801000 RDI: 00007f58d9b1e000
Jan 30 22:14:43 darkstar kernel: RBP: 00007f58b4dc79c0 R08: 0000000000000000 R09: 0000000000000001
Jan 30 22:14:43 darkstar kernel: R10: 00007f58d76000c0 R11: 0000000000000206 R12: 00007f5c029bb020
Jan 30 22:14:43 darkstar kernel: R13: 0000000002800000 R14: 00007f58b0dbd140 R15: 00007f58b0dbf700

```



smp_call_function_many

```

/**
 * smp_call_function_many(): Run a function on a set of other CPUs.
 * @mask: The set of cpus to run on (only runs on online subset).
 * @func: The function to run. This must be fast and non-blocking.
 * @info: An arbitrary pointer to pass to the function.
 * @wait: If true, wait (atomically) until function has completed
 *        on other CPUs.
 *
 * If @wait is true, then returns once @func has returned.
 *
 * You must not call this function with disabled interrupts or from a
 * hardware interrupt handler or from a bottom half handler. Preemption
 * must be disabled when calling this function.
 */

```



kernel/smp.c

循环等待

```
/* Send a message to all CPUs in the map */
arch_send_call_function_ipi_mask(cfd->cpumask_ipi);

if (wait) {
    for_each_cpu(cpu, cfd->cpumask) {
        call_single_data_t *csd;

        csd = per_cpu_ptr(cfd->csd, cpu);
        csd_lock_wait(csd);
    }
}
```

call_single_data_t



```
/**
 * smp_cond_load_relaxed() - (Spin) wait for cond with no ordering guarantees
 * @ptr: pointer to the variable to wait on
 * @cond: boolean expression to wait for
 *
 * Equivalent to using READ_ONCE() on the condition variable.
 *
 * Due to C lacking lambda expressions we load the value of *ptr into a
 * pre-named variable @VAL to be used in @cond.
 */
#ifndef smp_cond_load_relaxed
#define smp_cond_load_relaxed(ptr, cond_expr) ({ \
    typeof(ptr) __PTR = (ptr); \
    typeof(*ptr) VAL; \
    for (;;) { \
        VAL = READ_ONCE(*__PTR); \
        if (cond_expr) \
            break; \
        cpu_relax(); \
    } \
    VAL; \
})
#endif
```

忙等



flush_tlb_mm_range

arch > x86 > mm > C tlb.c > ...

```
728 void flush_tlb_mm_range(struct mm_struct *mm, unsigned long start,
729                          unsigned long end, unsigned int stride_shift,
730                          bool freed_tables)
731 {
732     int cpu;
```

```
754
755     if (mm == this_cpu_read(cpu_tlbstate.loaded_mm)) {
756         VM_WARN_ON(irqs_disabled());
757         local_irq_disable();
758         flush_tlb_func_local(&info, TLB_LOCAL_MM_SHOOTDOWN);
759         local_irq_enable();
760     }
761
762     if (cpumask_any_but(mm_cpumask(mm), cpu) < nr_cpu_ids)
763         flush_tlb_others(mm_cpumask(mm), &info);
764
765     put_cpu();
766 }
```

本地冲洗

通知伙伴

2020-01-30 20:26:00

Gruntz

Member From: Haskovo, Bulgaria

Registered: 2007-08-31 Posts: 285

Hello all,

I have a supemicro motherboard and two xeon x5650. I have 64GB of ram and several VMs on it.

I have one windows10, that i use for gaming (with gpu pass-through)

I have another, with archlinux for work (soft dev, testings... i use this one as my main desktop.GPU pass-through here too.)

I have 3-4-5 more, for database server, gitlab server, stuff like that.

From time to time my host machine crashes. It has a lot of "watchdog: BUG: soft lockup - CPU#4 stuck for 45s! [worker:131043]" messages.

Each time different processor. After it appears, the computer slowly becomes unresponsive, and eventually hangs completely.

Do you have any clue what this could be? I read about the problem. It appears when a cpu is stuck ot task for a long time or something like that.

But that is normal for VMs. Can I work around it?

Best regards,

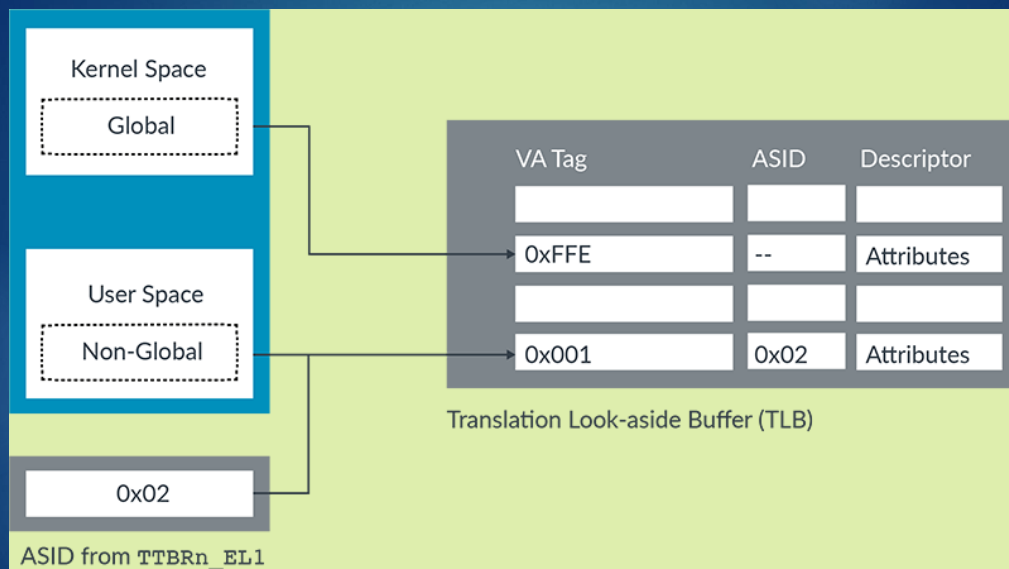
盛格聖

<https://bbs.archlinux.org/viewtopic.php?id=252523>

ASID

Address Space Identifiers

威格黎



威格黎

* <https://developer.arm.com/documentation/101811/0101/Address-spaces-in-AArch64>

Table K15-31 TLB maintenance system instructions

Register	Description, see
CFPRCTX	CFPRCTX
CPFRCTX	CPFRCTX
DTLBIALL	DTLBIALL
DTLBIASID	DTLBIASID
DTLBMVA	DTLBMVA
DVPRCTX	DVPRCTX
ITLBIALL	ITLBIALL
ITLBIASID	ITLBIASID
ITLBMVA	ITLBMVA
TLBIALL	TLBIALL
TLBIALLH	TLBIALLH
TLBIALLHIS	TLBIALLHIS
TLBIALLIS	TLBIALLIS
TLBIALLNSNH	TLBIALLNSNH
TLBIALLNSNHIS	TLBIALLNSNHIS
TLBIASID	TLBIASID
TLBIASIDIS	TLBIASIDIS
TLBIIPAS2	TLBIIPAS2
TLBIIPAS2IS	TLBIIPAS2IS
TLBIIPAS2L	TLBIIPAS2L
TLBIIPAS2LIS	TLBIIPAS2LIS
TLBMVA	TLBMVA
TLBMVAA	TLBMVAA
TLBMVAAIS	TLBMVAAIS
TLBMVAAL	TLBMVAAL
TLBMVAALIS	TLBMVAALIS
TLBMVAH	TLBMVAH
TLBMVAHIS	TLBMVAHIS
TLBMVAIS	TLBMVAIS
TLBMVAL	TLBMVAL

TLB maintenance system instructions

DTLBIASID, **Data** TLB Invalidate by ASID match

ITLBIASID, **Instruction** TLB Invalidate by ASID match

TLBIASID, TLB Invalidate by ASID match

TLBI <operation>{, <Xt>}

{R}<type><level><shareability>{NXS}



TLB maintenance instruction syntax

The A64 syntax for TLB maintenance instructions is:

TLBI <operation>{, <Xt>}

Where:

<operation> Is one of ALLE1{NXS}, ALLE2{NXS}, ALLE3{NXS}, ALLE1IS{NXS}, ALLE2IS{NXS}, ALLE3IS{NXS}, ALLE10S{NXS}, ALLE20S{NXS}, ALLE30S{NXS}, VMALLE1{NXS}, VMALLE1IS{NXS}, VMALLE10S{NXS}, VMALLS12E1{NXS}, VMALLS12E1IS{NXS}, VMALLS12E10S{NXS}, [ASIDE1{NXS}](#), [ASIDE1IS{NXS}](#), [ASIDE10S{NXS}](#), {R}VA{L}E1{NXS}, {R}VA{L}E2{NXS}, {R}VA{L}E3{NXS}, {R}VA{L}E1IS{NXS}, {R}VA{L}E2IS{NXS}, {R}VA{L}E3IS{NXS}, {R}VA{L}E10S{NXS}, {R}VA{L}E20S{NXS}, {R}VA{L}E30S{NXS}, {R}VAA{L}E1{NXS}, {R}VAA{L}E1IS{NXS}, {R}VAA{L}E10S{NXS}, {R}IPAS2{L}E1{NXS}, {R}IPAS2{L}E1IS{NXS}, or {R}IPAS2{L}E10S{NXS}.

<operation> has a structure of {R}<type><level><shareability>{NXS} where:

R When present, indicates that the function applies to all TLBs that are within a determined address range, see [TLB range maintenance instructions](#) on page D5-2828. When not present, indicates that the function applies to all TLBs at a single address that contain entries that could be used by the PE that executes the TLBI instruction.

<type> Is one of:

ALL All translations used at <level>.

For the scope of ALL instructions, see [ALL](#) on page D5-2824.

The ALL instructions are valid for all values of <level>.

VMALL All stage 1 translations used at <level> with the current VMID, if appropriate.



C tlbflush.h X

arch > arm64 > include > asm > C tlbflush.h

```


static inline void flush_tlb_kernel_range(unsigned long start, unsigned long end)
{
    unsigned long addr;

    if ((end - start) > (MAX_TLBI_OPS * PAGE_SIZE)) {
        flush_tlb_all();
        return;
    }

    start = __TLBI_VADDR(start, 0);
    end = __TLBI_VADDR(end, 0);

    dsb(ishst);
    for (addr = start; addr < end; addr += 1 << (PAGE_SHIFT - 12))
        __tlbi(vaale1is, addr);
    dsb(ish);
    isb();
}

```

TLB表项示例




```

arch > arm64 > mm > context.c
398 }
399 arch_initcall(asids_update_limit);
400
401 static int asids_init(void)
402 {
403     asid_bits = get_cpu_asid_bits();
404     atomic64_set(&asid_generation, ASID_FIRST_VERSION);
405     asid_map = kcalloc(BITS_TO_LONGS(NUM_USER_ASIDS), sizeof(*asid_map),
406                       GFP_KERNEL);
407     if (!asid_map)
408         panic("Failed to allocate bitmap for %lu ASIDs\n",
409               NUM_USER_ASIDS);
410
411     pinned_asid_map = kcalloc(BITS_TO_LONGS(NUM_USER_ASIDS),
412                              sizeof(*pinned_asid_map), GFP_KERNEL);
413     nr_pinned_asids = 0;
414
415     /*
416      * We cannot call set_reserved_asid_bits() here because CPU
417      * caps are not finalized yet, so it is safer to assume KPTI
418      * and reserve kernel ASID's from beginning.
419      */
420     if (IS_ENABLED(CONFIG_UNMAP_KERNEL_AT_EL0))
421         set_kpti_asid_bits(asid_map);
422     return 0;
423 }
424 early_initcall(asids_init);

```

≡ Kconfig X

arch > arm64 > ≡ Kconfig

715 config CAVIUM_ERRATUM_27456

config CAVIUM_ERRATUM_27456

bool "Cavium erratum 27456: Broadcast TLBI instructions may cause icache corruption"

default y

help

On ThunderX T88 pass 1.x through 2.1 parts, broadcast TLBI instructions may cause the icache to become corrupted if it contains data for a non-current ASID. The fix is to invalidate the icache when changing the mm context.

If unsure, say Y.

```
arch > x86 > mm > C tlb.c > ...
```


```

51  /*
52  * The x86 feature is called PCID (Process Context Identifier). It is similar
53  * to what is traditionally called ASID on the RISC processors.
54  *
55  * We don't use the traditional ASID implementation, where each process/mm gets
56  * its own ASID and flush/restart when we run out of ASID space.
57  *
58  * Instead we have a small per-cpu array of ASIDs and cache the last few mm's
59  * that came by on this CPU, allowing cheaper switch_mm between processes on
60  * this CPU.
61  *
62  * We end up with different spaces for different things. To avoid confusion we
63  * use different names for each of them:
64  *
65  * ASID - [0, TLB_NR_DYN_ASIDS-1]
66  *       the canonical identifier for an mm
67  *
68  * kPCID - [1, TLB_NR_DYN_ASIDS]
69  *         the value we write into the PCID part of CR3; corresponds to the
70  *         ASID+1, because PCID 0 is special.
71  *
72  * uPCID - [2048 + 1, 2048 + TLB_NR_DYN_ASIDS]
73  *         for KPTI each mm has two address spaces and thus needs two
74  *         PCID values, but we can still do with a single ASID denomination
75  *         for each mm. Corresponds to kPCID + 2048.

```

X86上叫
PCID，
代码里也
使用ASID





威格整

```
char *ptr = new char [16]; // memory colored
```



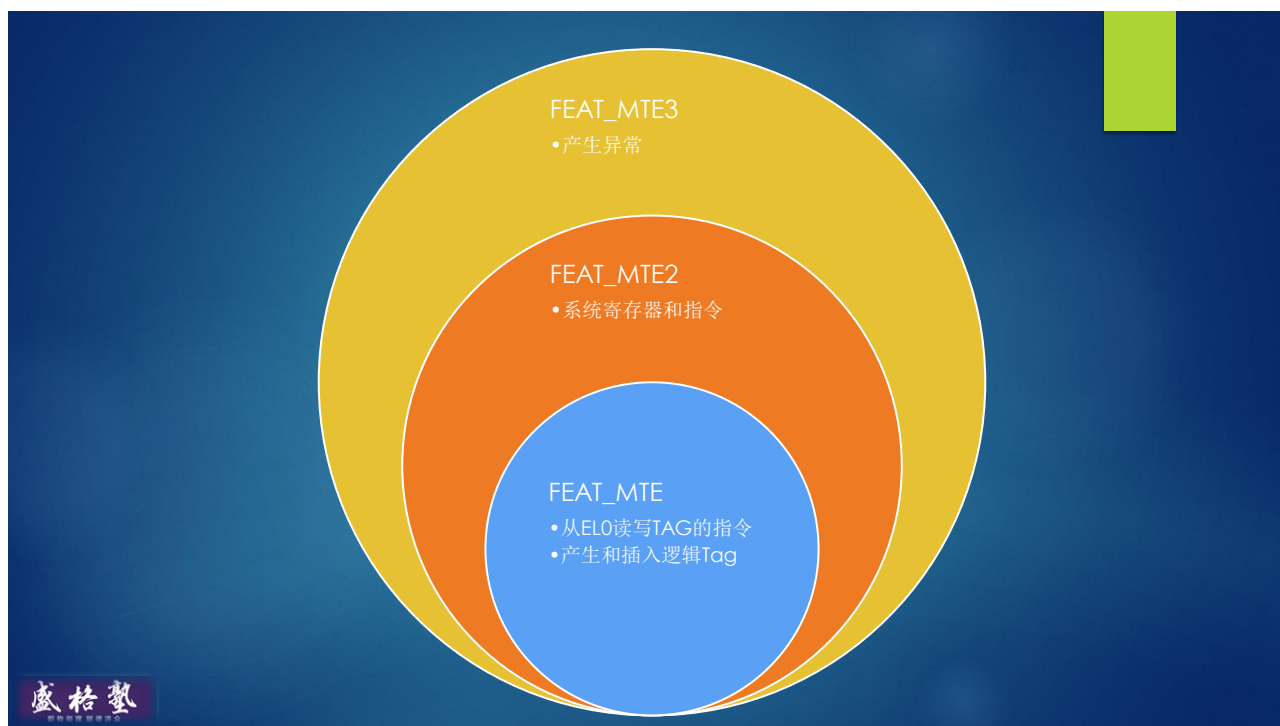
```
ptr[17] = 42; // color mismatch -> overflow
```

```
delete [] ptr; // memory re-coloured on free
```



```
ptr[10] = 10; // color mismatch -> use-after-free
```

威格整



→ security.googleblog.com/2019/08/adopting-arm-memory-tagging-extension.html

Adopting the Arm Memory Tagging Extension in Android

August 2, 2019

Posted by Kostya Serebryany, Google Core Systems, and Sudhi Herle, Android Security & Privacy Team

As part of our continuous commitment to improve the security of the Android ecosystem, we are partnering with Arm to design the [memory tagging extension](#) (MTE). [Memory safety bugs](#), common in C and C++, remain one of the largest vulnerabilities in the Android platform and although there have been previous [hardening efforts](#), memory safety bugs comprised more than half of the high priority security bugs in Android 9. Additionally, memory safety bugs manifest as hard to diagnose reliability problems, including sporadic crashes or silent data corruption. This reduces user satisfaction and [increases the cost of software development](#). Software testing tools, such as [ASAN](#) and [HWASAN](#) help, but their applicability on current hardware is limited due to noticeable overheads.

MTE, a hardware feature, aims to further mitigate these memory safety bugs by enabling us to detect them with low overhead. It has two execution modes:

<https://security.googleblog.com/2019/08/adopting-arm-memory-tagging-extension.html>

Chapter D6**Memory Tagging Extension**

D6.1	Introduction	D6-2840
D6.2	Allocation Tags	D6-2841
D6.3	Tag checking	D6-2842
D6.4	Tagged and Untagged Addresses	D6-2843
D6.5	PE access to Allocation Tags	D6-2844
D6.6	Enabling the Memory Tagging Extension	D6-2845
D6.7	PE handling of Tag Check Fault	D6-2846
D6.8	PE generation of Tag Checked and Tag Unchecked accesses	D6-2848

盛格塾
时格致 格致格



切问而近思

欢迎关注格友公众号

盛格塾
时格致 格致格