

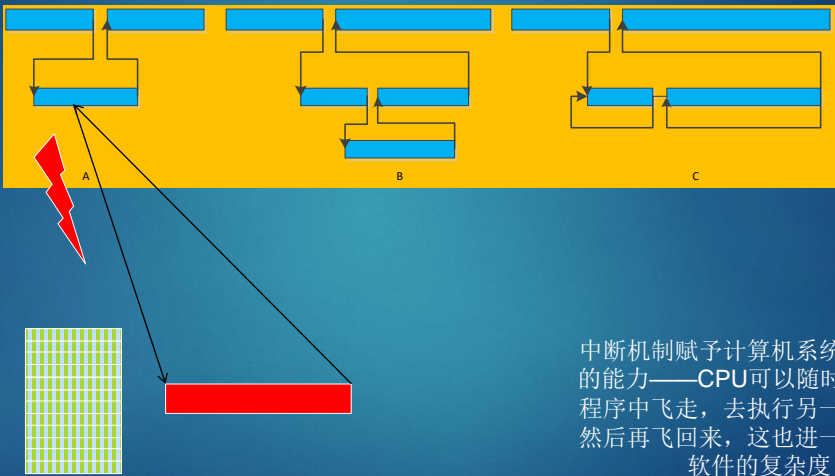


在调试器下理解ARMv8

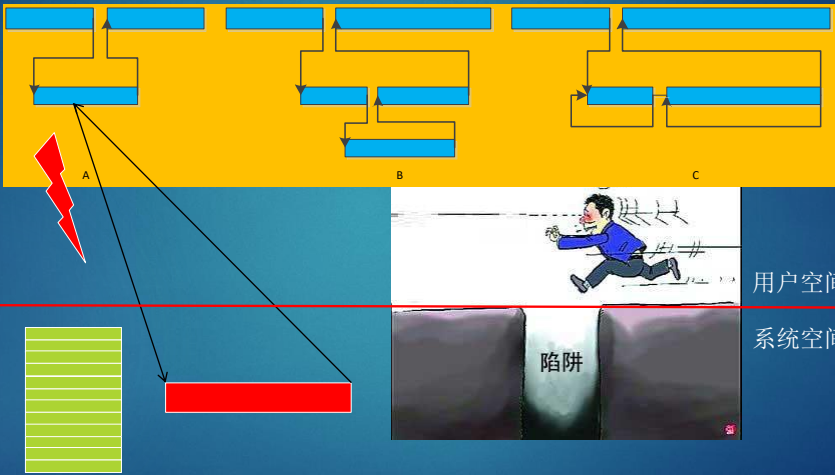
——异常和中断处理

张银奎 2022-1-16 中国·上海·古鹤坡塘





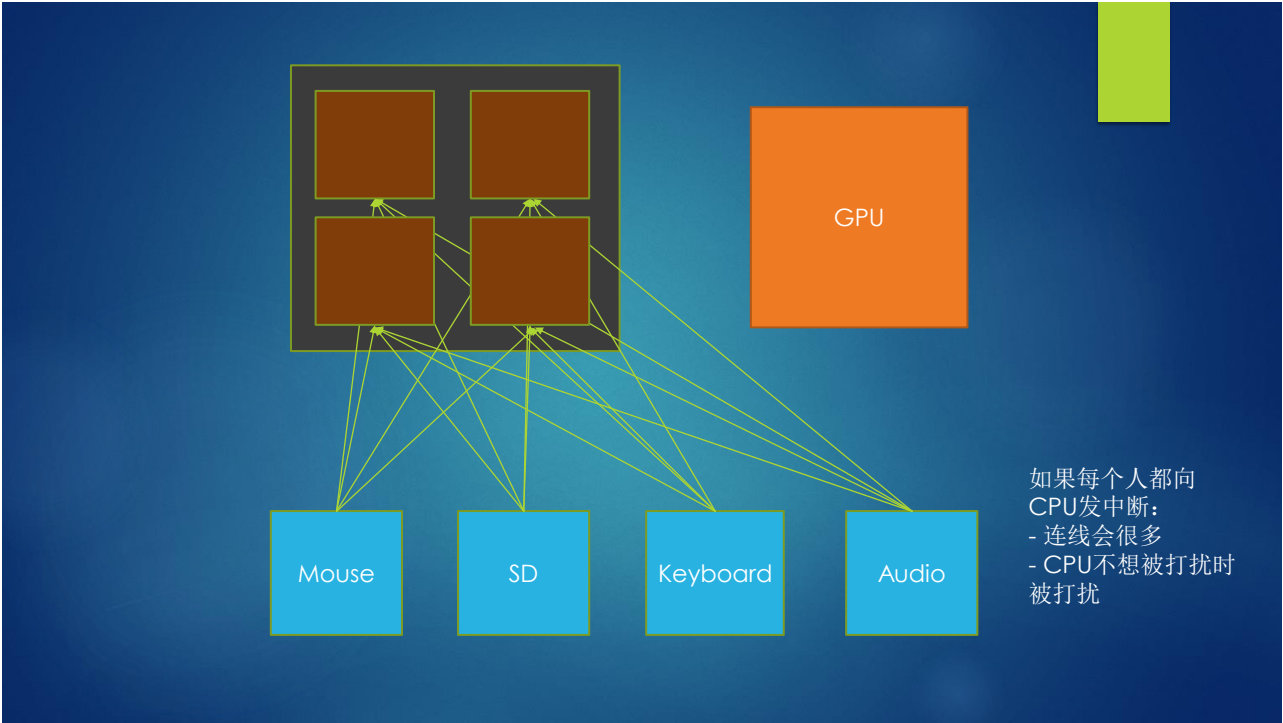
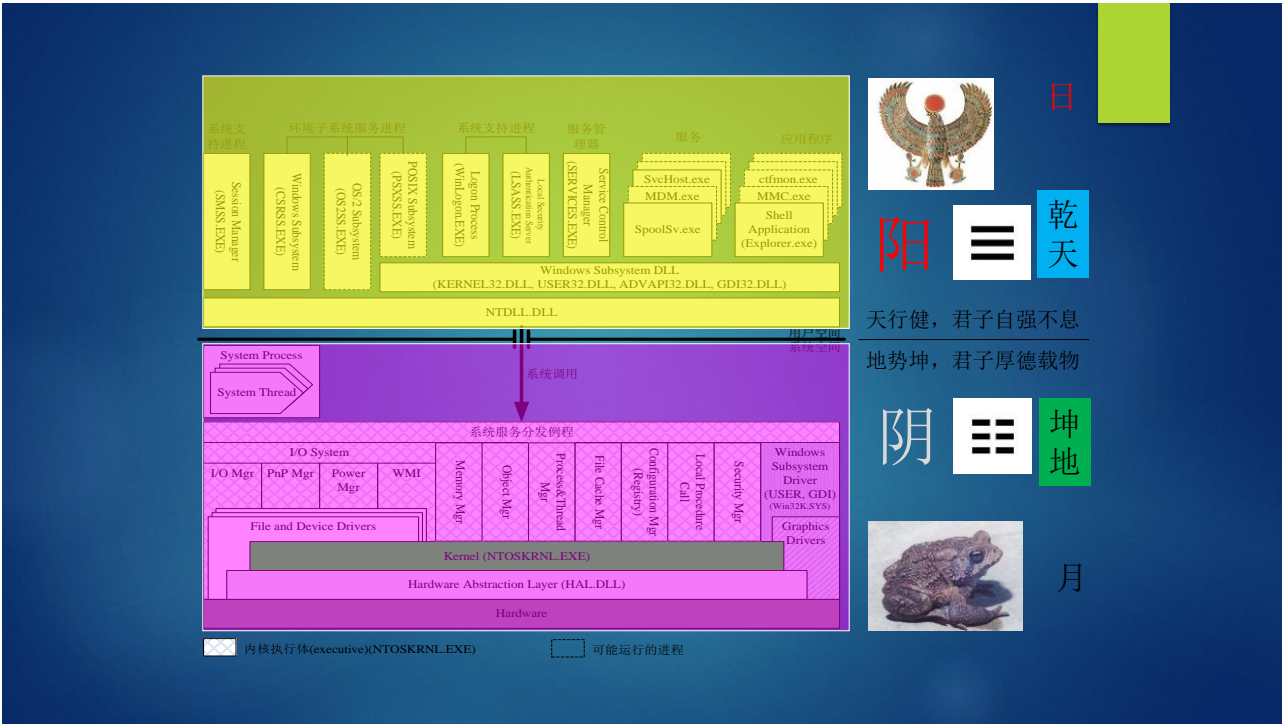
中断机制赋予计算机系统一种飞越的能力——CPU可以随时从当前的程序中飞走，去执行另一段程序，然后再飞回来，这也进一步增加了软件的复杂度

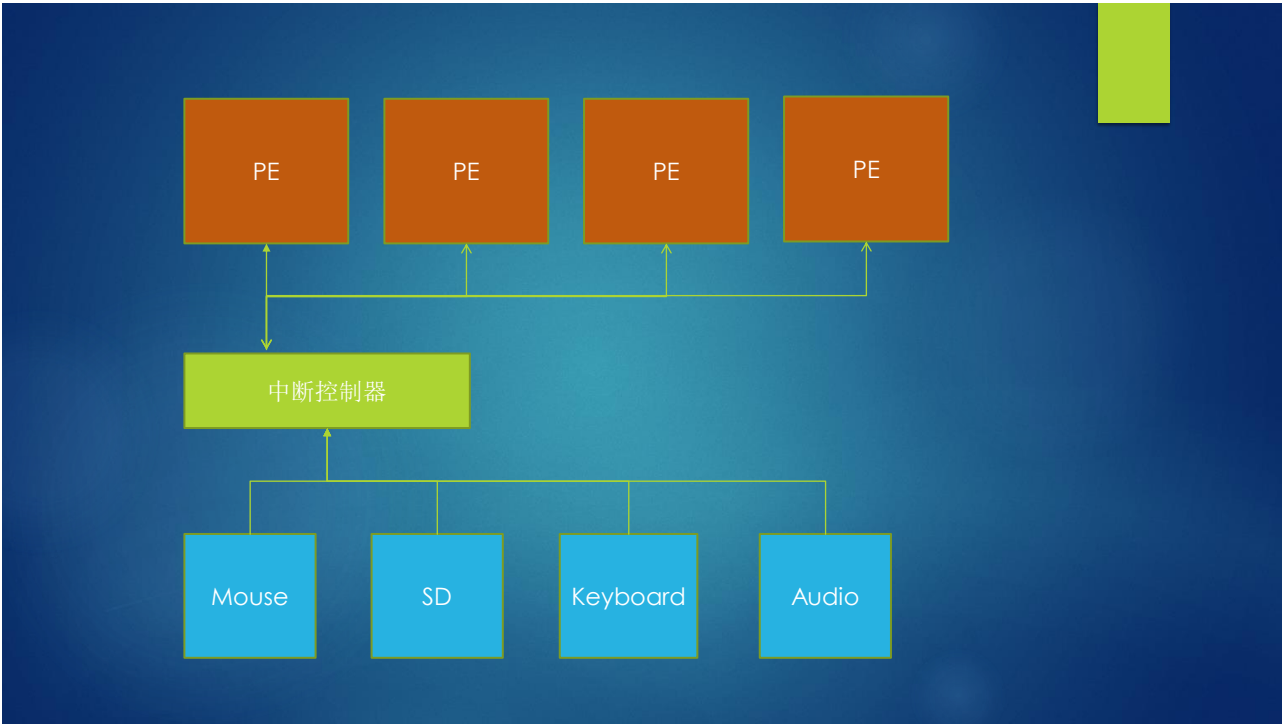


用户空间

系统空间

陷阱





CPU需要个“秘书”来管理中断

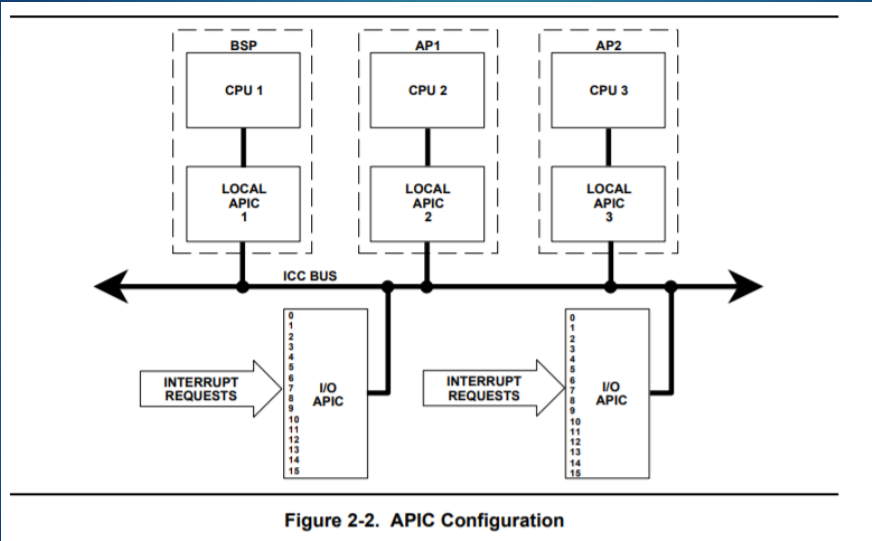
总线上应该有个“管家”来收纳和协调设备的中断

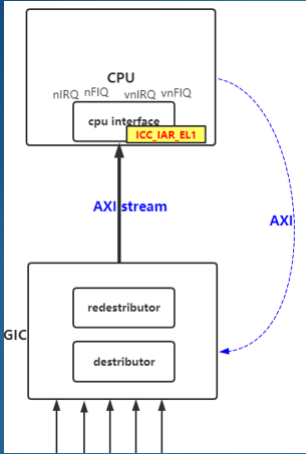
这个秘书便是中断控制器，X86上称为APIC，arm上称为GIC

C = Controller



X86上定义两种APIC





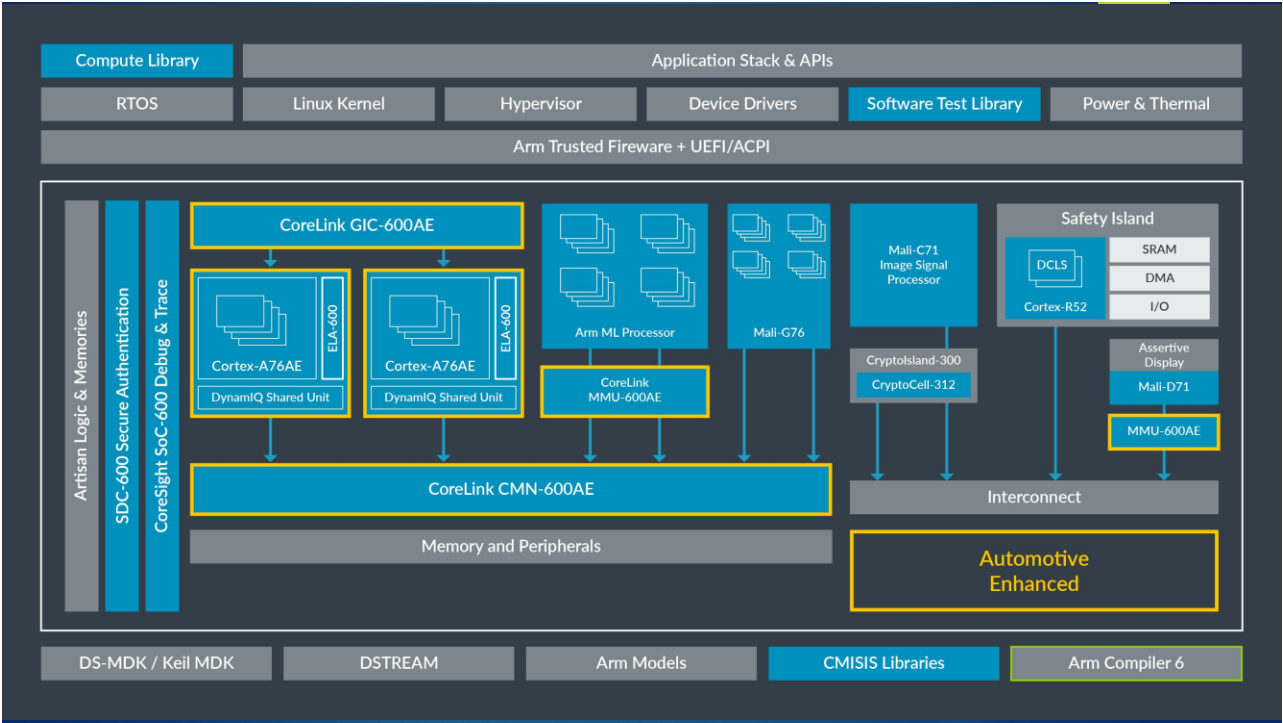
- ARM上没有明确定义两种APIC
- 把CPU上的叫CPU接口
 - 把总线上的叫GIC

A53 TRM中的GIC CPU Interface描述

- 9: Generic Interrupt Controller CPU Interface
 - 9.1 About the GIC CPU Interface
 - 9.2 GIC programmers model
 - 9.2.1 Memory map
 - 9.2.2 CPU interface register summary
 - 9.2.3 CPU interface register descriptions
 - 9.2.4 Virtual interface control register summary
 - 9.2.5 Virtual interface control register descriptions
 - 9.2.6 Virtual CPU interface register summary
 - 9.2.7 Virtual CPU interface register descriptions
- 10: Generic Timer

这些东西在X86上就是Local APIC

GIC = I/O APIC
GIC CPU Interface = Local APIC



Arm GIC架构规约

ARM® Generic Interrupt Controller

Architecture version 2.0

Architecture Specification

- ▶ 相当于x86的APIC手册
- ▶ 相对独立的GIC规约
- ▶ +214页

<https://developer.arm.com/documentation/ihl0048/latest>



GDK8中使用的A53实现的是GICv4

A53 TRM

9.1 About the GIC CPU Interface

The GIC CPU Interface, when integrated with an external distributor for supporting and managing interrupts in a cluster system. It provides:

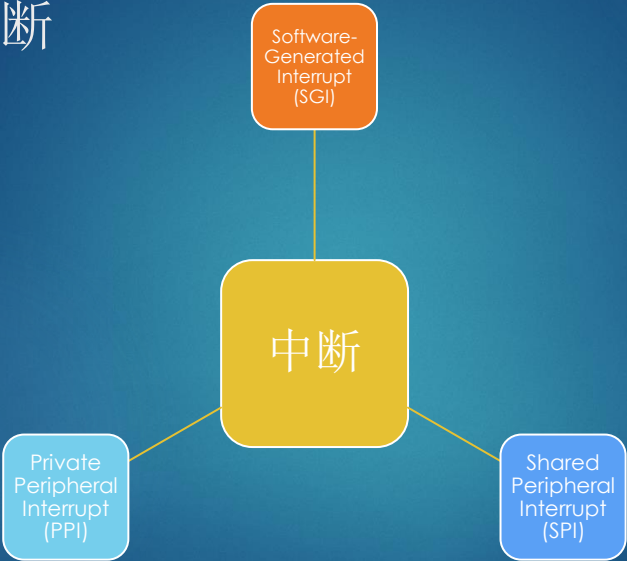
- Registers for managing:
 - Interrupt sources.
 - Interrupt behavior.
 - Interrupt routing to one or more cores.

The Cortex-A53 processor implements the GIC CPU interface as described in the Generic Interrupt Controller (GICv4) architecture. This interfaces with an external interrupt distributor component within the system.

The GICv4 architecture supports:

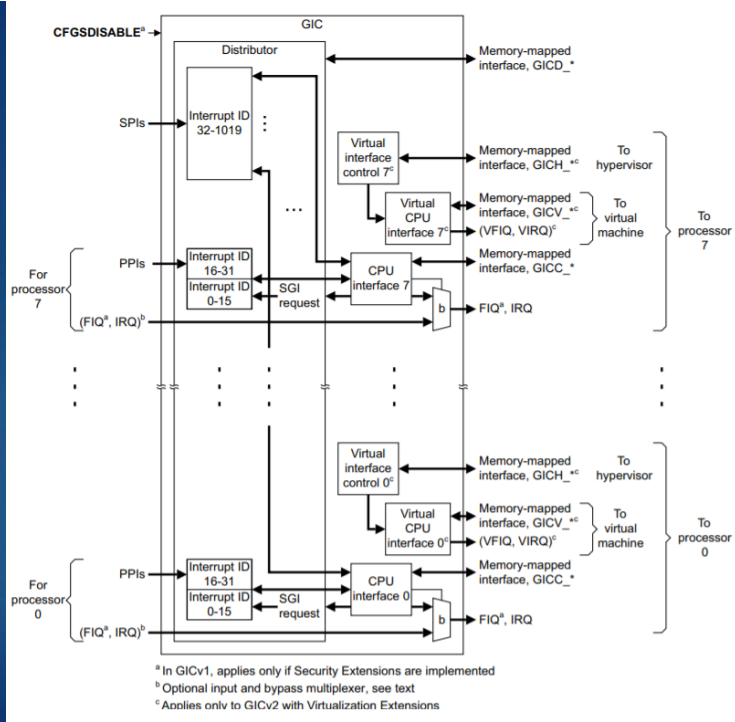
The Cortex-A53 processor implements the GIC CPU interface as described in the Generic Interrupt Controller (GICv4) architecture.

三类中断



中断ID分配

INTID	Interrupt Type	Notes
0 - 15	SGIs	Banked per PE
16 - 31	PPIs	Banked per PE
32 - 1019	SPIs	-
1020 - 1023	Special interrupt number	Used to signal special cases, see section 5.3
1024 - 8191	Reserved	-
8192 and greater	LPIs	The upper boundary is IMPLEMENTATION DEFINED



来自GIC架构规约



Chapter 7 Generic Interrupt Controller (GIC)

7.1 Overview

There is a generic interrupt controller(GIC400) in RK3328 which generates physical interrupts to Cortex-A53. It has two interfaces, the distributor interface connects to the interrupt source, and the CPU interface connects to Cortex-A53. The details of CPU interface connectivity are shown in the following table.

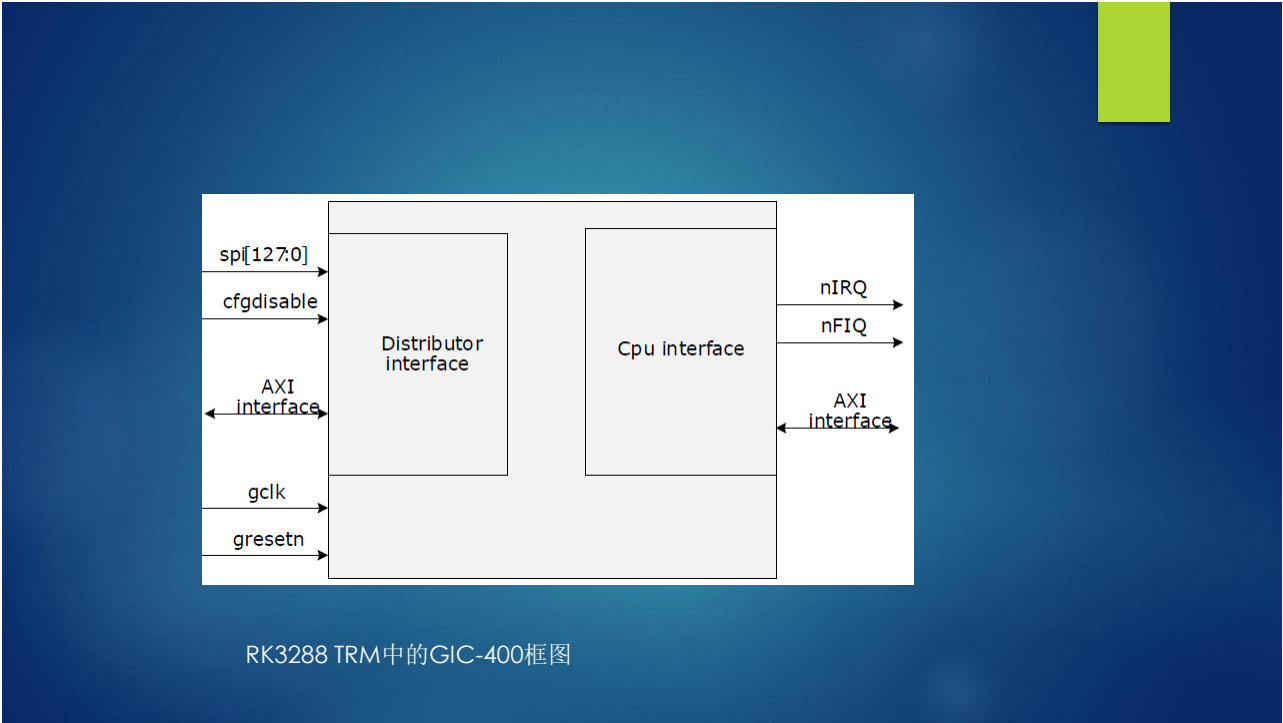
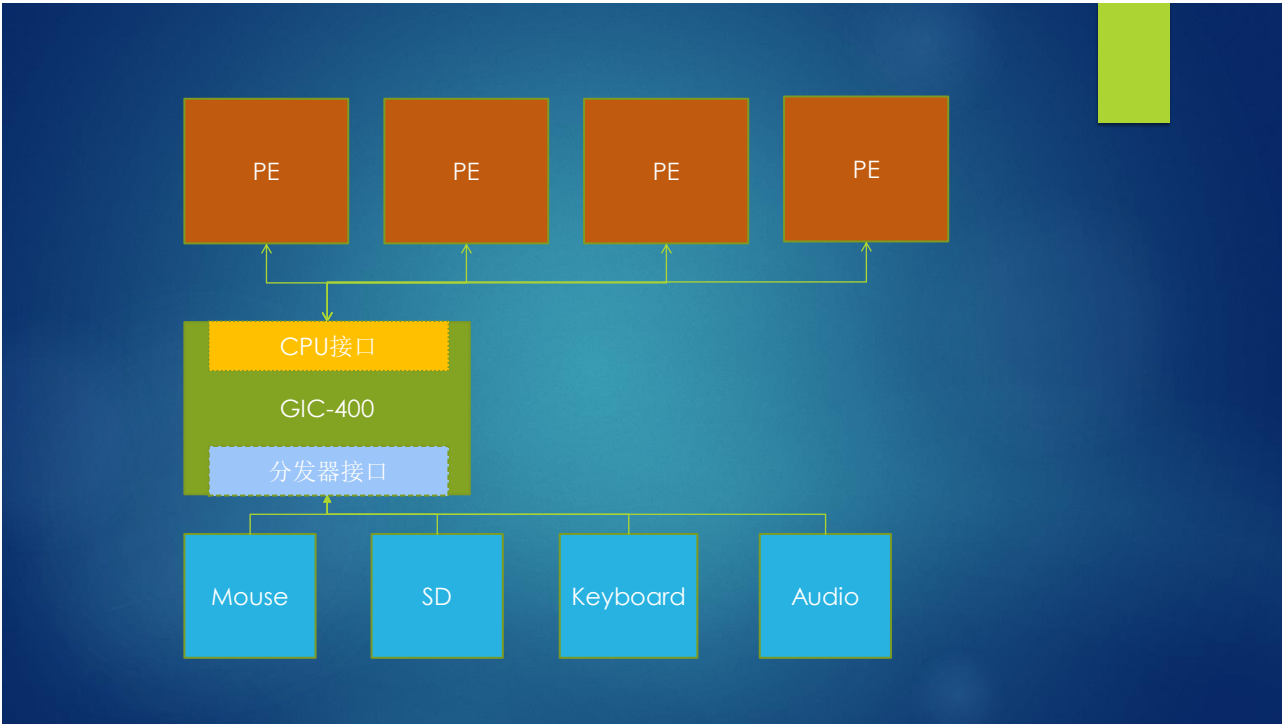
Table 1-1 CPU interface connectivity

CPU Interface Number	Connectivity
CPU interface 0	CPU0
CPU interface 1	CPU1
CPU interface 2	CPU2
CPU interface 3	CPU3

It supports the following features:

- Supports 128 hardware interrupt inputs
- Masking of any interrupts
- Prioritization of interrupts
- Distribution of the interrupts to the target Cortex-A53 processor(s)
- Generation of interrupts by software
- Supports Security Extensions

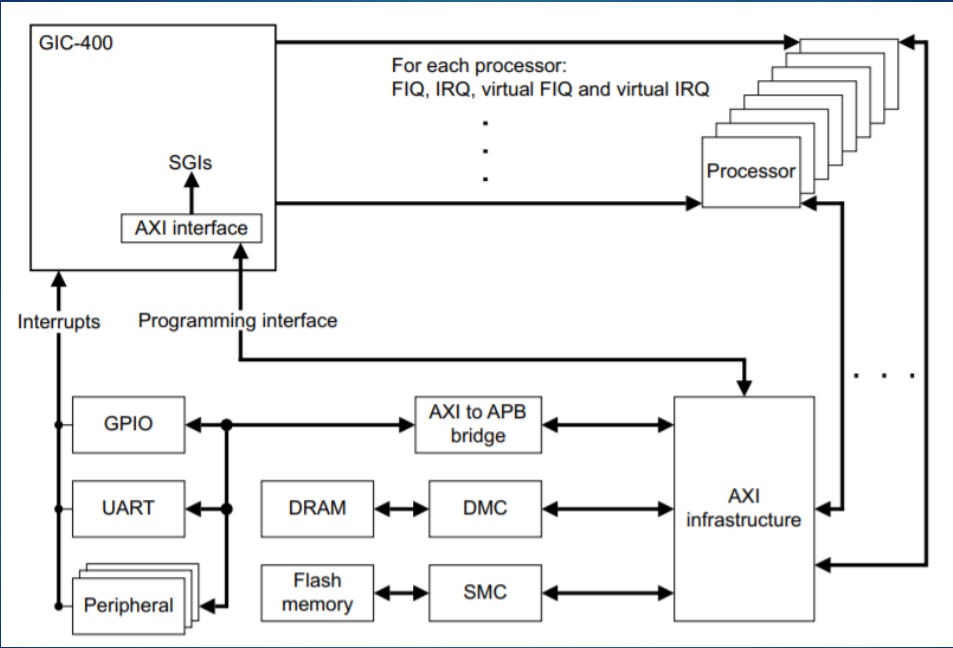
GDk8的
RK3328中
集成的
GIC是
GIC-400



CoreLink™ GIC-400 Generic Interrupt
Controller

Revision: r0p1

Technical Reference Manual



GIC-400 TRM中的工作原理图

```
interrupt-controller@ff811000 {  
    reg = <0x0 0xff811000 0x0 0x1000  
          0x0 0xff812000 0x0 0x2000  
          0x0 0xff814000 0x0 0x2000  
          0x0 0xff816000 0x0 0x2000>;  
    interrupts = <0x1 0x9 0xf04>;  
    compatible = "arm,gic-400";  
    #interrupt-cells = <0x3>;  
    #address-cells = <0x0>;  
    phandle = <0x1>;  
    interrupt-controller;  
};
```

寄存器接口描述，分别为：

/* GIC Dist */

/* GIC CPU */

/* GIC VCPU Control */

/* GIC VCPU */

在父GIC上的中断源，
也叫GIC
Maintenance IRQ

配置每个中断源的
参数个数，三个，
见下一页

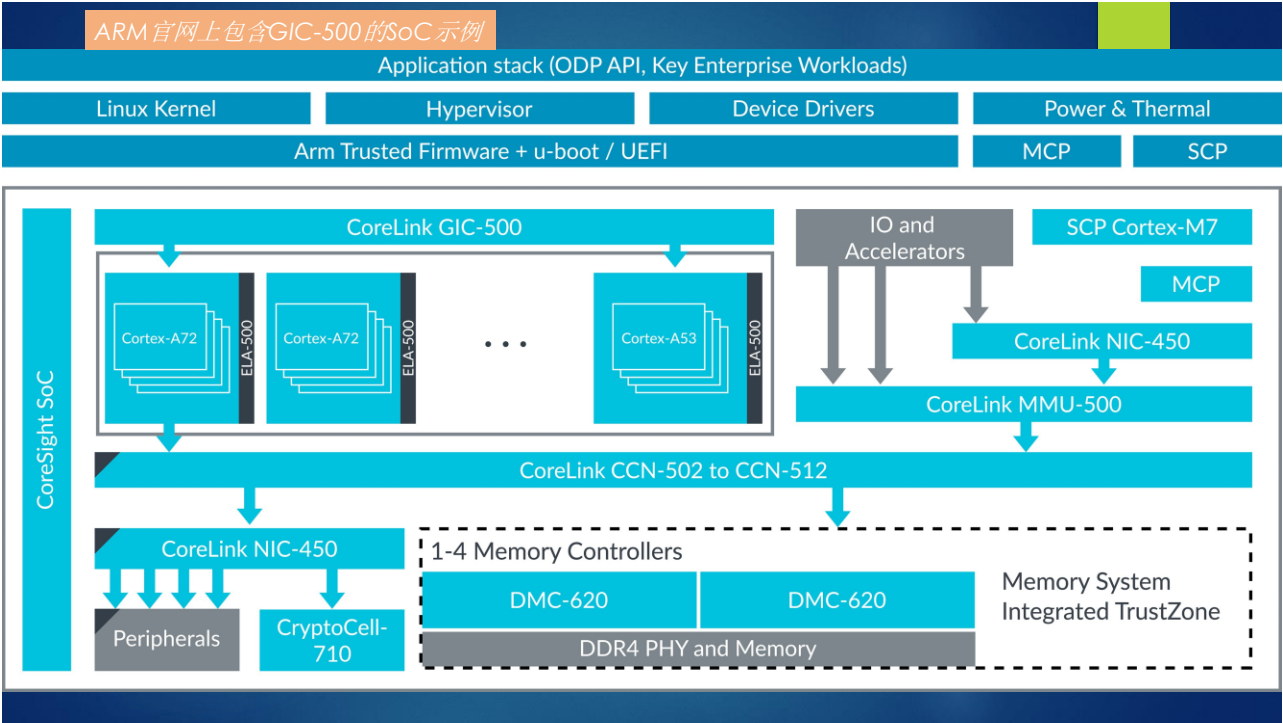
来自GDK8的DTS文件

GIC-400是对GIC架构规约的一种实现

GICv1-4是宏架构，给上层软件的接口，长期稳定，手册叫ARM GIC-400是微架构，考虑硬件成本和场景的一种实现，手册叫TRM

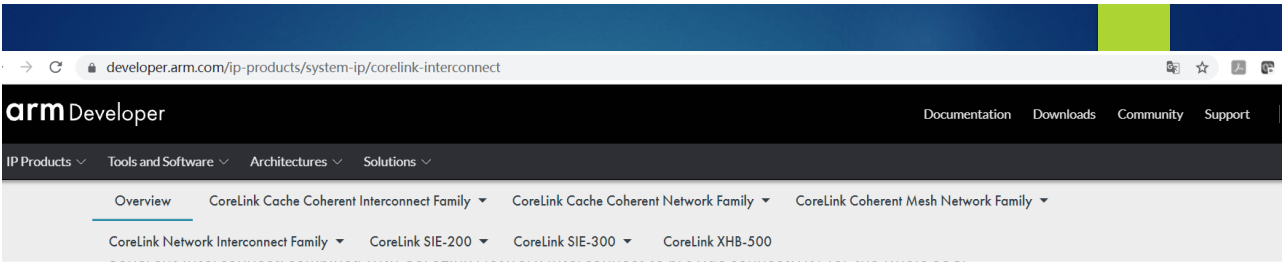


ARM官网上包含GIC-500的SoC示例

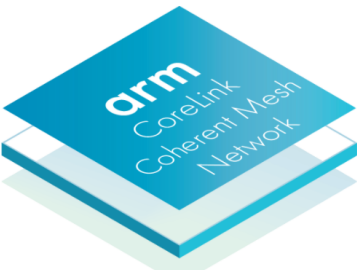


图中术语

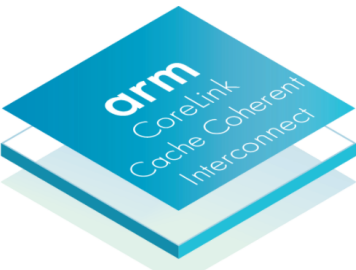
- ▶ CCN - **C**ache **C**oherent **N**etwork
- ▶ CoreLink – Core Link
- ▶ GIC - Generic Interrupt Controller
- ▶ ITS - Interrupt Translation Services



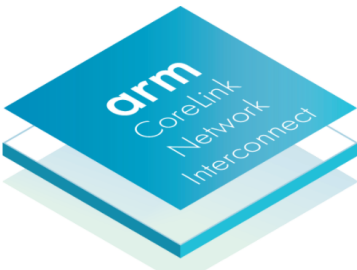
CoreLink Coherent Mesh Network



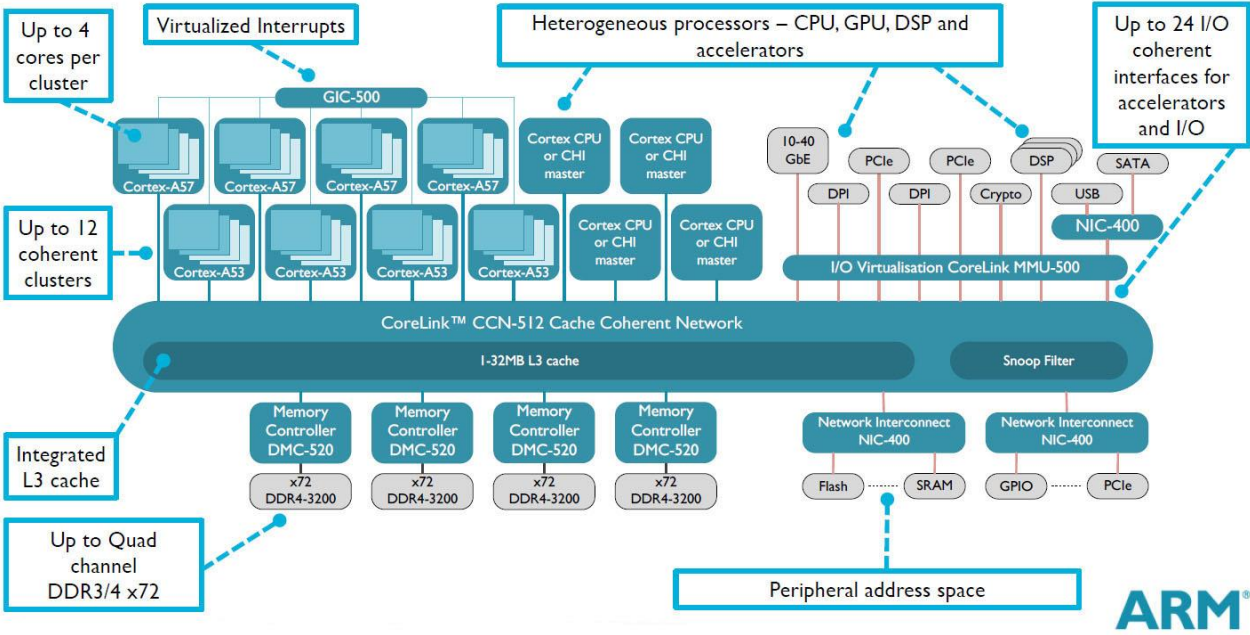
CoreLink Cache Coherent Interconnect



CoreLink Network Interconnect



ARM's CCN-512 Mixed Traffic Infrastructure SoC Framework



```
drivers > irqchip > C irq-gic-v3.c
1 x /*
```

- C irq-gic-common.c
- C irq-gic-common.h
- C irq-gic-pm.c
- C irq-gic-realview.c
- C irq-gic-v2m.c
- C irq-gic-v3-its-pci-msi.c
- C irq-gic-v3-its-platform-msi.c
- C irq-gic-v3-its.c
- C irq-gic-v3.c
- C irq-gic-v4.c
- C irq-gic.c

GIC的管理代码，或者说GIC本身的驱动

配置中断源

```
serial@ff110000 {
    reg = <0x0 0xff110000 0x0 0x0>;
    dmas = <0xc 0x2 0xc 0x2>;
    interrupts = <0x0 0x37 0x4>;
    pinctrl-0 = <0x35 0x36>;
    reg-shift = <0x2>;
    compatible = "rockchip,rk3328-uart";
    clock-names = "baudclk", "apb1_pclk";
    clocks = <0x2 0x26 0x2 0xd2>;
    status = "okay";
    reg-io-width = <0x4>;
    phandle = <0xa7>;
    pinctrl-names = "default";
};
```

中断类型，0代表SPI, 1代表PPI

4 = active high level-sensitive

中断源ID 37 38 39 40 41 42 43 44 45 46 47 48 49 50 51 52 53 54 55 56 57 58 59 60 61 62 63 64 65 66 67 68 69 70 71 72 73 74 75 76 77 78 79 80 81 82 83 84 85 86 87 88 89 90 91 92 93 94 95 96 97 98 99 100 101 102 103 104 105 106 107 108 109 110 111 112 113 114 115 116 117 118 119 120 121 122 123 124 125 126 127 128 129 130 131 132 133 134 135 136 137 138 139 140 141 142 143 144 145 146 147 148 149 150 151 152 153 154 155 156 157 158 159 160 161 162 163 164 165 166 167 168 169 170 171 172 173 174 175 176 177 178 179 180 181 182 183 184 185 186 187 188 189 190 191 192 193 194 195 196 197 198 199 200 201 202 203 204 205 206 207 208 209 210 211 212 213 214 215 216 217 218 219 220 221 222 223 224 225 226 227 228 229 230 231 232 233 234 235 236 237 238 239 240 241 242 243 244 245 246 247 248 249 250 251 252 253 254 255 256 257 258 259 260 261 262 263 264 265 266 267 268 269 270 271 272 273 274 275 276 277 278 279 280 281 282 283 284 285 286 287 288 289 290 291 292 293 294 295 296 297 298 299 300 301 302 303 304 305 306 307 308 309 310 311 312 313 314 315 316 317 318 319 320 321 322 323 324 325 326 327 328 329 330 331 332 333 334 335 336 337 338 339 340 341 342 343 344 345 346 347 348 349 350 351 352 353 354 355 356 357 358 359 360 361 362 363 364 365 366 367 368 369 370 371 372 373 374 375 376 377 378 379 380 381 382 383 384 385 386 387 388 389 390 391 392 393 394 395 396 397 398 399 400 401 402 403 404 405 406 407 408 409 410 411 412 413 414 415 416 417 418 419 420 421 422 423 424 425 426 427 428 429 430 431 432 433 434 435 436 437 438 439 440 441 442 443 444 445 446 447 448 449 450 451 452 453 454 455 456 457 458 459 460 461 462 463 464 465 466 467 468 469 470 471 472 473 474 475 476 477 478 479 480 481 482 483 484 485 486 487 488 489 490 491 492 493 494 495 496 497 498 499 500 501 502 503 504 505 506 507 508 509 510 511 512 513 514 515 516 517 518 519 520 521 522 523 524 525 526 527 528 529 530 531 532 533 534 535 536 537 538 539 540 541 542 543 544 545 546 547 548 549 550 551 552 553 554 555 556 557 558 559 560 561 562 563 564 565 566 567 568 569 570 571 572 573 574 575 576 577 578 579 580 581 582 583 584 585 586 587 588 589 590 591 592 593 594 595 596 597 598 599 600 601 602 603 604 605 606 607 608 609 610 611 612 613 614 615 616 617 618 619 620 621 622 623 624 625 626 627 628 629 630 631 632 633 634 635 636 637 638 639 640 641 642 643 644 645 646 647 648 649 650 651 652 653 654 655 656 657 658 659 660 661 662 663 664 665 666 667 668 669 670 671 672 673 674 675 676 677 678 679 680 681 682 683 684 685 686 687 688 689 690 691 692 693 694 695 696 697 698 699 700 701 702 703 704 705 706 707 708 709 710 711 712 713 714 715 716 717 718 719 720 721 722 723 724 725 726 727 728 729 730 731 732 733 734 735 736 737 738 739 740 741 742 743 744 745 746 747 748 749 750 751 752 753 754 755 756 757 758 759 760 761 762 763 764 765 766 767 768 769 770 771 772 773 774 775 776 777 778 779 780 781 782 783 784 785 786 787 788 789 790 791 792 793 794 795 796 797 798 799 800 801 802 803 804 805 806 807 808 809 810 811 812 813 814 815 816 817 818 819 820 821 822 823 824 825 826 827 828 829 830 831 832 833 834 835 836 837 838 839 840 841 842 843 844 845 846 847 848 849 850 851 852 853 854 855 856 857 858 859 860 861 862 863 864 865 866 867 868 869 870 871 872 873 874 875 876 877 878 879 880 881 882 883 884 885 886 887 888 889 890 891 892 893 894 895 896 897 898 899 900 901 902 903 904 905 906 907 908 909 910 911 912 913 914 915 916 917 918 919 920 921 922 923 924 925 926 927 928 929 930 931 932 933 934 935 936 937 938 939 940 941 942 943 944 945 946 947 948 949 950 951 952 953 954 955 956 957 958 959 960 961 962 963 964 965 966 967 968 969 970 971 972 973 974 975 976 977 978 979 980 981 982 983 984 985 986 987 988 989 990 991 992 993 994 995 996 997 998 999

以GPIO形式配置中断源

```
wireless-bluetooth {
    pinctrl-0 = <0x99>;
    pinctrl-1 = <0x9a>;
    uart_rts_gpios = <0x6e 0xa 0x1>;
    compatible = "bluetooth-platdata";
    clock-names = "ext_clock";
    clocks = <0x2 0x1e>;
    BT,wake_gpio = <0x6e 0x17 0x0>;
    BT,reset_gpio = <0x6e 0x15 0x0>;
    status = "okay";
    pinctrl-names = "default", "rts_gpio";
    BT,wake_host_irq = <0x6e 0x1a 0x0>;
};
```



蓝牙和WIFI无线模块

蓝牙模块通过GPIO管脚连到SoC, 触发中断

33:	0	0	0	0	GICv2	53	Level	eth0
42:	4950	0	0	0	GICv2	64	Level	ff370000.vop
105:	0	0	0	0	gpio1	26	Edge	bt default wake host_irq
144:	57485	0	0	0	gpio3	1	Level	bcm5dh_sdmhc
175:	0	0	0	0	GICv2	39	Level	ff360000.rkvdec

```
geduer@gdk8:~$ cat /proc/interrupts
CPU0          CPU1          CPU2          CPU3
 1:           0             0             0             0      GICv2 29 Edge      arch_timer
 2:       132012       140116       97041       182397      GICv2 30 Edge      arch_timer
 5:         931             0             0             0      GICv2 32 Level     ff1f0000.dmac
 6:           0             0             0             0      GICv2 33 Level     ff1f0000.dmac
 9:           0             0             0             0      GICv2 48 Level     ehci_hcd:usb2
10:           0             0             0             0      GICv2 49 Level     ohci_hcd:usb3
11:           0             0             0             0      GICv2 90 Level     rockchip_thermal
12:           0             0             0             0      GICv2 41 Level     ff350000.vpu_service, ff351000.avsd
14:           0       34067             0             0      GICv2 82 Level     rk_pwm_irq
16:         4128             0             0             0      GICv2 67 Level     ff3c0000.hdmi, dw-hdmi-cec
18:           0             0             0             0      GICv2 44 Level     dw-mci
19:        64883             0             0             0      GICv2 46 Level     dw-mci
20:         1190             0             0             0      GICv2 87 Level     serial
21:       277893             0             0             0      GICv2 36 Level     dw-mci
22:           0             0             0             0      GICv2 122 Level    Mali_GP
23:           0             0             0             0      GICv2 119 Level    Mali_GP_MMU
24:           0             0             0             0      GICv2 125 Level    Mali_PP_Broadcast
25:           0             0             0             0      GICv2 120 Level    Mali_PP0
26:           0             0             0             0      GICv2 121 Level    Mali_PP0_MMU
27:           0             0             0             0      GICv2 123 Level    Mali_PP1
28:           0             0             0             0      GICv2 124 Level    Mali_PP1_MMU
29:           0             0             0             0      GICv2 112 Level    ff280000.saradc
30:           0             0             0             0      GICv2 115 Level    ff430000.hdmiphy
31:           0             0             0             0      GICv2 127 Level    ff330000.h265e
32:           0             0             0             0      GICv2 159 Edge     debug-signal
33:           0             0             0             0      GICv2 53 Level     eth0
42:        4950             0             0             0      GICv2 64 Level     ff370000.vop
105:          0             0             0             0      gpio1 26 Edge     bt_default_wake_host_irq
144:       33864             0             0             0      gpio3 1 Level     bcmsdh_sdmmc
175:           0             0             0             0      GICv2 39 Level     ff360000.rkvdec
176:           0             0             0             0      GICv2 129 Level    ff340000.vepu
177:           0             0             0             0      GICv2 100 Level    rockchip_u3phy
178:          2             0             0             0      GICv2 55 Level     ff580000.usb, ff580000.usb, dwc2_hstgt:usb1
179:           0             0             0             0      GICv2 89 Level     debug
180:           0             0             0             0      GICv2 91 Level     rockchip_usb2phy_bvalid
181:           0             0             0             0      GICv2 93 Level     rockchip_usb2phy
182:           0             0             0             0      GICv2 92 Level     rockchip_usb2phy_id
183:           0             0             0             0      GICv2 94 Level     rockchip_usb2phy
184:        2709             0             0             0      GICv2 99 Level     xhci-hcd:usb4
```

[0.000000] GIC: Using split EOI/Deactivate mode


```
[ 0.000000] NR_IRQS:64 nr_irqs:64 0
[ 0.180409] genirq: Setting trigger mode 8 for irq 32 failed (gic_set_type+0x0/0x64)
[ 0.528883] ffl10000.serial: ttyS0 at MMIO 0xff110000 (irq = 20, base_baud = 1500000) is a 16550A
[ 0.819068] dwc2 ff580000.usb: irq 178, io mem 0xff580000
[ 1.471936] ehci-platform ff5c0000.usb: irq 9, io mem 0xff5c0000
[ 1.488440] ohci-platform ff5d0000.usb: irq 10, io mem 0xff5d0000
[ 1.553017] xhci-hcd xhci-hcd.9.auto: irq 184, io mem 0xff600000
[ 1.623229] dwmmc_rockchip ff500000.dwmmc: DW MMC controller at irq 18,32 bit host data width,256 deep fifo
[ 1.650944] dwmmc_rockchip ff520000.dwmmc: DW MMC controller at irq 19,32 bit host data width,256 deep fifo
[ 2.175684] dwmmc_rockchip ff5f0000.dwmmc: DW MMC controller at irq 21,32 bit host data width,256 deep fifo
[ 2.233655] of_get_named_gpiod_flags: parsed 'WIFI,host_wake_irq' property of node '/wireless-wlan[0]' - status (0)
[ 2.233663] [WLAN_RFKILL]: wlan_platdata_parse_dt: get property: WIFI,host_wake_irq = 97, flags = 0.
[ 2.238390] of_get_named_gpiod_flags: parsed 'BT,wake_host_irq' property of node '/wireless-bluetooth[0]' - status (0)
[ 2.238394] [BT_RFKILL]: bluetooth_platdata_parse_dt: get property: BT,wake_host_irq = 58.
[ 2.238478] [BT_RFKILL]: Request irq for bt wakeup host
[ 2.238520] [BT_RFKILL]: ** disable irq and enable wake
[ 2.238528] [BT_RFKILL]: ** irq wake (105) is enabled on GPIO 58
[ 7.526901] [WLAN_RFKILL]: rockchip_wifi_get_oob_irq: Enter
[ 7.526912] dhd_wlan_init_gpio: WL_HOST_WAKE=-1, oob_irq=144, oob_irq_flags=0x414
[ 7.881570] bcm5dh_oob_intr_unregister: irq is not registered
[ 8.671179] bcm5dh_oob_intr_register: HW_OOB irq=144 flags=0x4
```

晃动GDK8 USB 3.0口上的鼠标，这个数字会变化

此处的4是流水号，代表是系统里的第4个USB设备

```
184:      2709      0      0      0      GICv2  99 Level  xhci-hcd:usb4
```

```
[ 2.008771] xhci-hcd xhci-hcd.9.auto: irq 184, io mem 0xff600000
```


从DTS中获取配置

```
irq = platform_get_irq_byname(dwc3_pdev, "dwc_usb3");
irq = platform_get_irq(dwc3_pdev, 0);
```

```
dwc->xhci_resources[1].start = irq;
dwc->xhci_resources[1].end = irq;
dwc->xhci_resources[1].flags = res->flags;
dwc->xhci_resources[1].name = res->name;
```

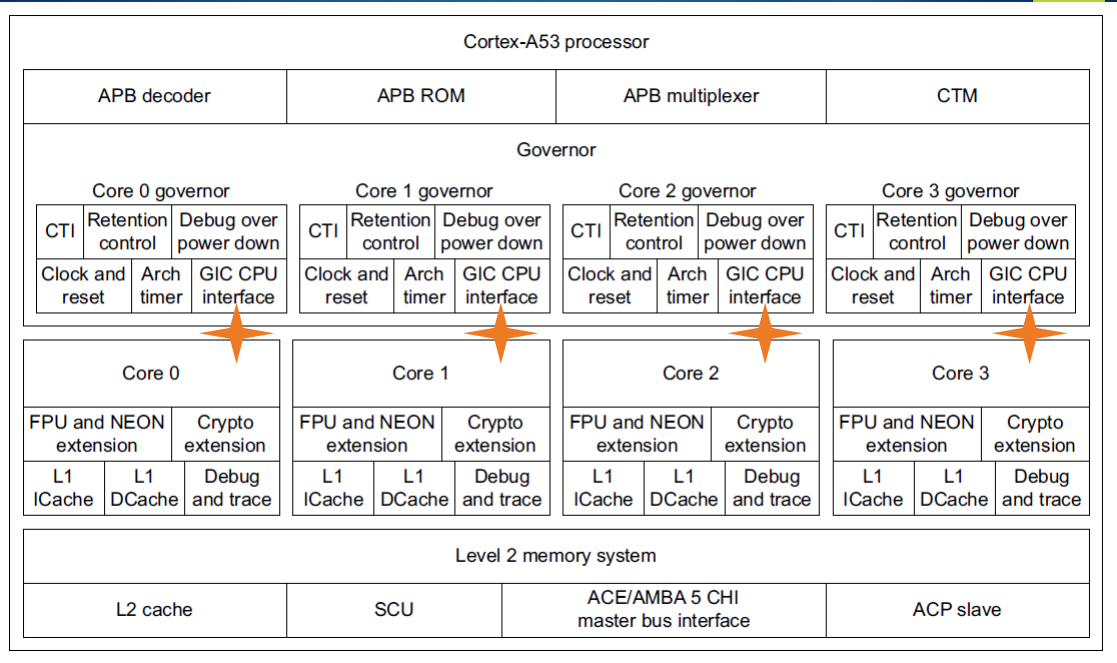
```
ret = platform_device_add_resources(xhci, dwc->xhci_resources,
DWC3_XHCI_RESOURCES_NUM);
```

```
ret = platform_device_add(xhci);
```

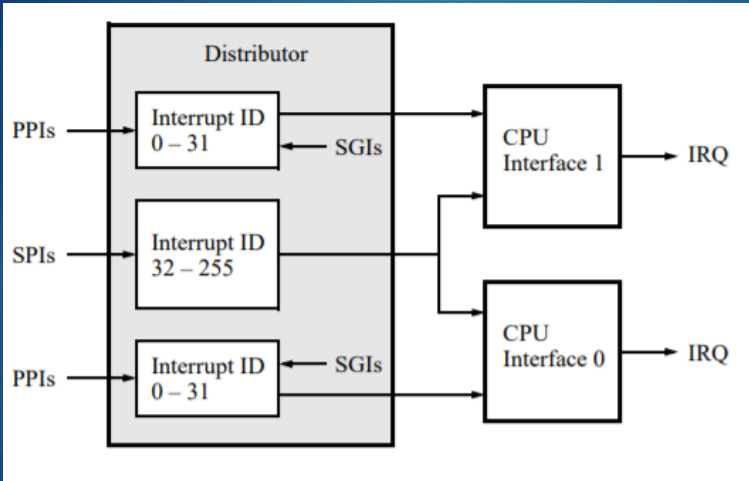
```
drivers > usb > dwc3 > C host.c > dwc3_host_init(dwc3 *)
22
23 int dwc3_host_init(struct dwc3 *dwc)
24 {
```

```
arch > arm64 > boot > dts > arm > foundation-v8.dts
102     };
103
104     smb {
105         compatible = "arm,vexpress,v2m-p1", "simple-bus";
106         arm,v2m-memory-map = "rs1";
107         #address-cells = <2>; /* SMB chipselect number and offset */
108         #size-cells = <1>;
109
110         ranges = <0 0 0 0x08000000 0x04000000>,
111             <1 0 0 0x14000000 0x04000000>,
112             <2 0 0 0x18000000 0x04000000>,
113             <3 0 0 0x1c000000 0x04000000>,
114             <4 0 0 0x0c000000 0x04000000>,
115             <5 0 0 0x10000000 0x04000000>;
116
117         #interrupt-cells = <1>;
118         interrupt-map-mask = <0 0 0 63>;
119         interrupt-map = <0 0 0 &gic 0 0 4>,
120             <0 0 1 &gic 0 1 4>,
121             <0 0 2 &gic 0 2 4>,
122             <0 0 3 &gic 0 3 4>,
123             <0 0 4 &gic 0 4 4>,
124             <0 0 5 &gic 0 5 4>,
125             <0 0 6 &gic 0 6 4>,
126             <0 0 7 &gic 0 7 4>,
127             <0 0 8 &gic 0 8 4>,
128             <0 0 9 &gic 0 9 4>,
129             <0 0 10 &gic 0 10 4>,
130             <0 0 11 &gic 0 11 4>;
```

中断映射，可以通过DTS定义规则，也可以在操作系统中设置中断亲缘性，Linux内核也有irqbalance服务（GDK8中没有启用）



CPU接口



PPI - private peripherals interrupts
SPI - Shared peripherals interrupts

https://ftp.intel.com/Public/Pub/fpgaup/pub/Teaching_Materials/current/Tutorials/Using_GIC.pdf

Offset	Name	Type	Reset	Description
0x0000	GICC_CTLR	RW	0x00000000	CPU Interface Control Register
0x0004	GICC_PMR	RW	0x00000000	Interrupt Priority Mask Register
0x0008	GICC_BPR	RW	0x00000002 (S) ^a 0x00000003 (NS) ^b	Binary Point Register
0x000C	GICC_IAR	RO	-	Interrupt Acknowledge Register
0x0010	GICC_EOIR	WO	-	End Of Interrupt Register
0x0014	GICC_RPR	RO	0x000000FF	Running Priority Register
0x0018	GICC_HPPIR	RO	0x000003FF	Highest Priority Pending Interrupt Register
0x001C	GICC_ABPR	RW	0x00000003	Aliased Binary Point Register
0x0020	GICC_AIAR	RO	-	Aliased Interrupt Acknowledge Register
0x0024	GICC_AEOIR	WO	-	Aliased End of Interrupt Register
0x0028	GICC_AHPPIR	RO	0x000003FF	Aliased Highest Priority Pending Interrupt Register
0x00D0	GICC_APR0	RW	0x00000000	<i>Active Priority Register on page 9-6</i>
0x00E0	GICC_NSAPR0	RW	0x00000000	Non-secure Active Priority Register
0x00FC	GICC_IIDR	RO	0x0034443B	<i>CPU Interface Identification Register on page 9-7</i>
0x1000	GICC_DIR	WO	-	Deactivate Interrupt Register

CPU接口寄存器

来自A53 TRM

两种访问方式



通过内存映射访问

Address	31	...	10	9	8	7	...	1	0	Register name
0xFFFFEC100	Unused								E	ICCICR
0xFFFFEC104	Unused						Priority			ICCPMR
0xFFFFEC10C	Unused						Interrupt ID			ICCIAR
0xFFFFEC110	Unused						Interrupt ID			ICCEOIR

- ▶ CPU通过虚拟地址访问，x86中称为Local APIC，每个CPU有自己的一份实例

4.4.2 Interrupt Priority Mask Register, GICC_PMR

The GICC_PMR characteristics are:

Purpose	Provides an interrupt priority filter. Only interrupts with higher priority than the value in this register are signaled to the processor.
----------------	--

– **Note**

Higher priority corresponds to a lower Priority field value.

Usage constraints If the GIC implements the Security Extensions then:

- a Non-secure access to this register can only read or write a value that corresponds to the lower half of the priority range, see [Interrupt grouping and interrupt prioritization on page 3-53](#).
- if a Secure write has programmed the GICC_PMR with a value that corresponds to a value in the upper half of the priority range then:
 - any Non-secure read of the GICC_PMR returns 0x00, regardless of the value held in the register
 - any Non-secure write to the GICC_PMR is ignored.

For more information see *Non-secure access to register fields for Group 0 interrupt priorities* on page 4-81.

When determining interrupt preemption, the priority value can be split into two parts, using the Binary Point register, [GICC_BPR](#).

Configurations	This register is available in all configurations of the GIC. If the GIC implements the Security Extensions, this register is Common.
-----------------------	--

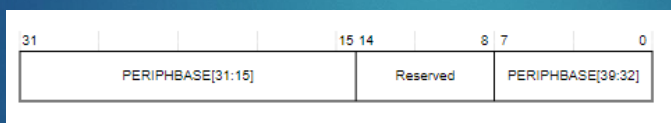
Attributes See the register summary in Table 4-2 on page 4-76.

Figure 4-25 shows the GICC_PMR bit assignments.



CBAR_EL1

- ▶ GIC内存映射寄存器的基地址寄存器
- ▶ Configuration Base Address Register
- ▶ EL1或者更高级别才能访问



ARM Cortex-A15 MPCore Processor Technical Reference Manual



中断和异常向量表

Offset from VBAR_EL1	Exception type	Exception set level
+0x000	Synchronous	Current EL with SP0
+0x080	IRQ/vIRQ	“
+0x100	FIQ/vFIQ	“
+0x180	SError/vSError	“
+0x200	Synchronous	Current EL with SPx
+0x280	IRQ/vIRQ	“
+0x300	FIQ/vFIQ	“
+0x380	SError/vSError	“
+0x400	Synchronous	Lower EL using ARM64
+0x480	IRQ/vIRQ	“
+0x500	FIQ/vFIQ	“
+0x580	SError/vSError	“
+0x600	Synchronous	Lower EL with ARM32
+0x680	IRQ/vIRQ	“
+0x700	FIQ/vFIQ	“
+0x780	SError/vSError	“

arch/arm64/kernel/entry.S

```

1      ENTRY(vectors)
2      ventry el1_sync_invalid // Synchronous EL1t
3      ventry el1_irq_invalid // IRQ EL1t
4      ventry el1_fiq_invalid // FIQ EL1t
5      ventry el1_error_invalid // Error EL1t
6
7      ventry el1_sync // Synchronous EL1h
8      ventry el1_irq // IRQ EL1h
9      ventry el1_fiq_invalid // FIQ EL1h
10     ventry el1_error_invalid // Error EL1h
11
12     ventry el0_sync // Synchronous 64-bit EL0
13     ventry el0_irq // IRQ 64-bit EL0
14     ventry el0_fiq_invalid // FIQ 64-bit EL0
15     ventry el0_error_invalid // Error 64-bit EL0
16
17     #ifdef CONFIG_COMPAT
18     ventry el0_sync_compat // Synchronous 32-bit EL0
19     ventry el0_irq_compat // IRQ 32-bit EL0
20     ventry el0_fiq_invalid_compat // FIQ 32-bit EL0
21     ventry el0_error_invalid_compat // Error 32-bit EL0
22     #else
23     ventry el0_sync_invalid // Synchronous 32-bit EL0
24     ventry el0_irq_invalid // IRQ 32-bit EL0
25     ventry el0_fiq_invalid // FIQ 32-bit EL0
26     ventry el0_error_invalid // Error 32-bit EL0
27     #endif
28     END(vectors)

```

每个ventry占32个字节，8条指令

* <https://eastrivervillage.com/Anatomy-of-Linux-system-call-in-ARM64/>

加载向量表

```

adr_l  x8, vectors          // load VBAR_EL1 with virtual
msr    vbar_el1, x8         // vector table address
lsb    // instruction sync barrier

```

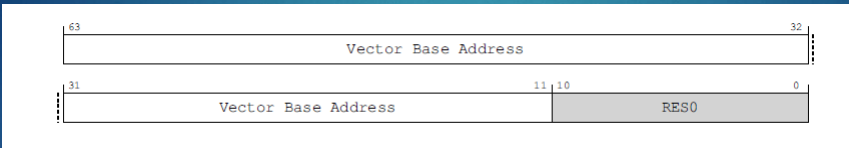
```
void ge_arm_sysregs(void)
{
    rd_arm_reg(ID_AA64ISAR0_EL1);
    rd_arm_reg(ID_AA64ISAR1_EL1);
    rd_arm_reg(ID_AA64MMFR0_EL1);
    rd_arm_reg(ID_AA64MMFR1_EL1);
    rd_arm_reg(ID_AA64PFR0_EL1);
    rd_arm_reg(ID_AA64PFR1_EL1);
    rd_arm_reg(ID_AA64DFR0_EL1);
    rd_arm_reg(ID_AA64DFR1_EL1);

    rd_arm_reg(MIDR_EL1);
    rd_arm_reg(MPIDR_EL1);
    rd_arm_reg(REVIDR_EL1);
    rd_arm_reg(VBAR_EL1);
    /* Unexposed register access causes SIGILL */
    rd_arm_reg(ID_MMFR0_EL1);
}
```



VBAR

Vector Base Address Register



Bits	Name	Function
[31:5]	VBA	Bits[31:5] of the base address of the low exception vectors.
[4:0]	-	Reserved.

AArch64 System Register Descriptions
D13.2 General system control registers

D13.2.140 VBAR_EL1, Vector Base Address Register (EL1)

```

1  el0_sync:
2  kernel_entry 0
3  mrs x25, esr_el1 // 把异常症状寄存器读到x25
4  lsr x24, x25, #ESR_ELx_EC_SHIFT // 把x25左移26位，得到的值（exception class）放入x24
5  cmp x24, #ESR_ELx_EC_SVC64 // SVC in 64-bit state
6  b.eq el0_svc
7  cmp x24, #ESR_ELx_EC_DABT_LOW // data abort in EL0
8  b.eq el0_da
9  cmp x24, #ESR_ELx_EC_IABT_LOW // instruction abort in EL0
10 b.eq el0_ia
11 cmp x24, #ESR_ELx_EC_FP_ASIMD // FP/ASIMD access
12 b.eq el0_fpsimd_acc
13 cmp x24, #ESR_ELx_EC_FP_EXC64 // FP/ASIMD exception
14 b.eq el0_fpsimd_exc
15 cmp x24, #ESR_ELx_EC_SYS64 // configurable trap
16 b.eq el0_sys
17 cmp x24, #ESR_ELx_EC_SP_ALIGN // stack alignment exception
18 b.eq el0_sp_pc
19 cmp x24, #ESR_ELx_EC_PC_ALIGN // pc alignment exception
20 b.eq el0_sp_pc
21 cmp x24, #ESR_ELx_EC_UNKNOWN // unknown exception in EL0
22 b.eq el0_undef
23 cmp x24, #ESR_ELx_EC_BREAKPT_LOW // debug exception in EL0
24 b.ge el0_dbg
25 b el0_inv

```

// arch/arm64/kernel/entry.S

```

[27948.980226] ID_AA64ISAR0_EL1 : 0x0000000000001120
[27948.980232] ID_AA64ISAR1_EL1 : 0x0000000000000000
[27948.980237] ID_AA64MMFR0_EL1 : 0x0000000000001122
[27948.980243] ID_AA64MMFR1_EL1 : 0x0000000000000000
[27948.980248] ID_AA64PFR0_EL1 : 0x0000000000000222
[27948.980254] ID_AA64PFR1_EL1 : 0x0000000000000000
[27948.980259] ID_AA64DFR0_EL1 : 0x0000000010305106
[27948.980265] ID_AA64DFR1_EL1 : 0x0000000000000000
[27948.980270] MIDR_EL1 : 0x00000000410fd034
[27948.980275] MPIDR_EL1 : 0x00000000080000003
[27948.980281] REVIDR_EL1 : 0x00000000000000180
[27948.980286] VBAR_EL1 : 0xfffff8008081800
[27948.980292] ID_MMFR0_EL1 : 0x0000000010201105

```

ESR_EL1

Exception Syndrome Register

Bits	Name	Function
[31:26]	EC	Exception Class: 0b100000 Instruction Abort that caused entry from a lower Exception level in AArch32 or AArch64. 0b100001 Instruction Abort that caused entry from a current Exception level in AArch64.
[25]	IL	Instruction Length for synchronous exceptions.
[24:10]	-	Reserved, RES0.
[9]	EA	External abort type. This bit indicates whether an AXI decode or slave error caused an abort. The possible values are: 0 External abort marked as DECERR. 1 External abort marked as SLVERR. For aborts other than external aborts this bit always returns 0.
[8]	-	Reserved, RES0.
[7]	S1PTW	When 1, indicates the instruction fault came from a second stage fault during a first stage translation table walk.
[6]	-	Reserved, RES0.
[5:0]	IFSC	Instruction Fault Status Code This field indicates the type of exception generated. The possible values are:

```
#define ESR_ELx_EC_UNKNOWN (0x00)
#define ESR_ELx_EC_WFx      (0x01)
/* Unallocated EC: 0x02 */
#define ESR_ELx_EC_CP15_32  (0x03)
#define ESR_ELx_EC_CP15_64  (0x04)
#define ESR_ELx_EC_CP14_MR  (0x05)
#define ESR_ELx_EC_CP14_LS  (0x06)
#define ESR_ELx_EC_FP_ASIMD (0x07)
#define ESR_ELx_EC_CP10_ID   (0x08) /* EL2 only */
#define ESR_ELx_EC_PAC       (0x09) /* EL2 and above */
/* Unallocated EC: 0x0A - 0x0B */
#define ESR_ELx_EC_CP14_64   (0x0C)
/* Unallocated EC: 0x0d */
#define ESR_ELx_EC_IILL      (0x0E)
/* Unallocated EC: 0x0F - 0x10 */
#define ESR_ELx_EC_SVC32     (0x11)
#define ESR_ELx_EC_HVC32     (0x12) /* EL2 only */
#define ESR_ELx_EC_SMC32     (0x13) /* EL2 and above */
/* Unallocated EC: 0x14 */
#define ESR_ELx_EC_SVC64     (0x15)
#define ESR_ELx_EC_HVC64     (0x16) /* EL2 and above */
#define ESR_ELx_EC_SMC64     (0x17) /* EL2 and above */
#define ESR_ELx_EC_SYS64     (0x18)
#define ESR_ELx_EC_SVE       (0x19)
```



```
#define ESR_ELx_EC_IMP_DEF      (0x1f)  /* EL3 only */
#define ESR_ELx_EC_IABT_LOW     (0x20)
#define ESR_ELx_EC_IABT_CUR     (0x21)
#define ESR_ELx_EC_PC_ALIGN     (0x22)
/* Unallocated EC: 0x23 */
#define ESR_ELx_EC_DABT_LOW     (0x24)
#define ESR_ELx_EC_DABT_CUR     (0x25)
#define ESR_ELx_EC_SP_ALIGN     (0x26)
/* Unallocated EC: 0x27 */
#define ESR_ELx_EC_FP_EXC32     (0x28)
/* Unallocated EC: 0x29 - 0x2B */
#define ESR_ELx_EC_FP_EXC64     (0x2C)
/* Unallocated EC: 0x2D - 0x2E */
#define ESR_ELx_EC_ERROR        (0x2F)
#define ESR_ELx_EC_BREAKPT_LOW  (0x30)
#define ESR_ELx_EC_BREAKPT_CUR  (0x31)
#define ESR_ELx_EC_SOFTSTP_LOW  (0x32)
#define ESR_ELx_EC_SOFTSTP_CUR  (0x33)
#define ESR_ELx_EC_WATCHPT_LOW  (0x34)
#define ESR_ELx_EC_WATCHPT_CUR  (0x35)
#define ESR_ELx_EC_BKPT32       (0x38)
#define ESR_ELx_EC_VECTOR32     (0x3A)  /* EL2 only */
#define ESR_ELx_EC_BRK64        (0x3C)
#define ESR_ELx_EC_MAX          (0x3F)
```

\arch\arm64\include\asm\esr.h

```
#define ESR_ELx_EC_SHIFT      (26)
#define ESR_ELx_EC_SVC64      (0x15)
```

系统服务入口

```

1  /*
2  * SVC handler.
3  */
4  .align 6
5  el0_svc:
6  adrp stbl, sys_call_table // load syscall table pointer
7  uxtw scno, w8 // syscall number in w8
8  mov sc_nr, #__NR_syscalls
9  el0_svc_naked: // compat entry point
10 stp x0, scno, [sp, #S_ORIG_X0] // save the original x0 and syscall number enable_dbg_and_irq
11 ct_user_exit 1
12
13 ldr x16, [tsk, #TI_FLAGS] // check for syscall hooks
14 tst x16, #_TIF_SYSCALL_WORK
15 b.ne __sys_trace
16 cmp scno, sc_nr // check upper syscall limit
17 b.hs ni_sys
18 ldr x16, [stbl, scno, lsl #3] // 将服务号左移3位，得到偏移，加上基地址，得到address in the syscall table
19 blr x16 // call sys_* routine
20 b ret_fast_syscall
21 ni_sys:
22 mov x0, sp
23 bl do_ni_syscall
24 b ret_fast_syscall
25 ENDPROC(el0_svc)

```

寄存器别名

```

1  /*
2  * These are the registers used in the syscall handler, and allow us to
3  * have in theory up to 7 arguments to a function - x0 to x6.
4  *
5  * x7 is reserved for the system call number in 32-bit mode.
6  */
7  sc_nr .req x25 // number of system calls
8  scno .req x26 // syscall number
9  stbl .req x27 // syscall table pointer
10 tsk .req x28 // current thread_info

```


系统服务表

```

1  #undef __SYSCALL
2  #define __SYSCALL(nr, sym) [nr] = sym,
3
4  /*
5   * The sys_call_table array must be 4K aligned to be accessible from
6   * kernel/entry.S.
7   */
8  void * const sys_call_table[__NR_syscalls] __aligned(4096) = {
9  [0 ... __NR_syscalls - 1] = sys_ni_syscall,
10 #include <asm/unistd.h>
11 };

```

arch/arm64/kernel/sys.c

调用系统服务

```

1  | 0x7fb7f407a8 <__GI___libc_write> stp x29, x30, [sp, #-48]!
2  | 0x7fb7f407ac <__GI___libc_write+4> adrp x3, 0x7fb7fd1000 <__libc_pthread_functions+184>
3  | 0x7fb7f407b0 <__GI___libc_write+8> mov x29, sp
4  | 0x7fb7f407b4 <__GI___libc_write+12> str x19, [sp, #16]
5  | 0x7fb7f407b8 <__GI___libc_write+16> sxtw x19, w0
6  | 0x7fb7f407bc <__GI___libc_write+20> ldr w0, [x3, #264]
7  | 0x7fb7f407c0 <__GI___libc_write+24> cbnz w0, 0x7fb7f407f0 <__GI___libc_write+72>
8  | 0x7fb7f407c4 <__GI___libc_write+28> mov x0, x19
9  | 0x7fb7f407c8 <__GI___libc_write+32> mov x8, #0x40 // #64
10 | 0x7fb7f407cc <__GI___libc_write+36>

```

SVC

SVC指令

SuperVisor Call.

Syntax

```
SVC{ cond } # imm
```

where:

cond

is an optional condition code.

imm

is an expression evaluating to an integer in the range:

- 0 to $2^{24}-1$ (a 24-bit value) in an ARM instruction.
- 0-255 (an 8-bit value) in a Thumb instruction.

https://www.keil.com/support/man/docs/armasm/armasm_dom1361289909139.htm

- ▶ 0b000000 Address size fault in TTBR0 or TTBR1.
- ▶ 0b000101 Translation fault, 1st level.
- ▶ 0b000110 Translation fault, 2nd level.
- ▶ 0b000111 Translation fault, 3rd level.
- ▶ 0b001001 Access flag fault, 1st level.
- ▶ 0b001010 Access flag fault, 2nd level.
- ▶ 0b001011 Access flag fault, 3rd level.
- ▶ 0b001101 Permission fault, 1st level.
- ▶ 0b001110 Permission fault, 2nd level.
- ▶ 0b001111 Permission fault, 3rd level.
- ▶ 0b010000 Synchronous external abort.
- ▶ 0b011000 Synchronous parity error on memory access.
- ▶ 0b010101 Synchronous external abort on translation table walk, 1st level.
- ▶ 0b010110 Synchronous external abort on translation table walk, 2nd level.
- ▶ 0b010111 Synchronous external abort on translation table walk, 3rd level.
- ▶ 0b011101 Synchronous parity error on memory access on translation table walk, 1st level.
- ▶ 0b011110 Synchronous parity error on memory access on translation table walk, 2nd level.
- ▶ 0b011111 Synchronous parity error on memory access on translation table walk, 3rd level.
- ▶ 0b100001 Alignment fault.
- ▶ 0b100010 Debug event.

Instruction Fault Status Code

切问而近思

欢迎关注格友公众号

